Transformer Model (2/2): From Shallow to Deep

Shusen Wang

Multi-Head Attention

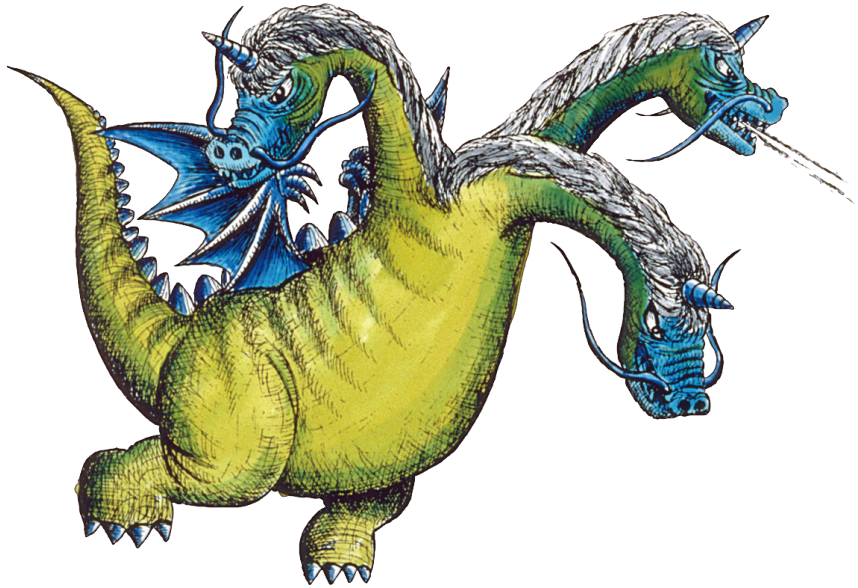

Single-Head Self-Attention

- Self-attention layer: $\mathbf{C} = \mathrm{Attn}(\mathbf{X}, \mathbf{X})$.
- This is called "single-head self-attention".

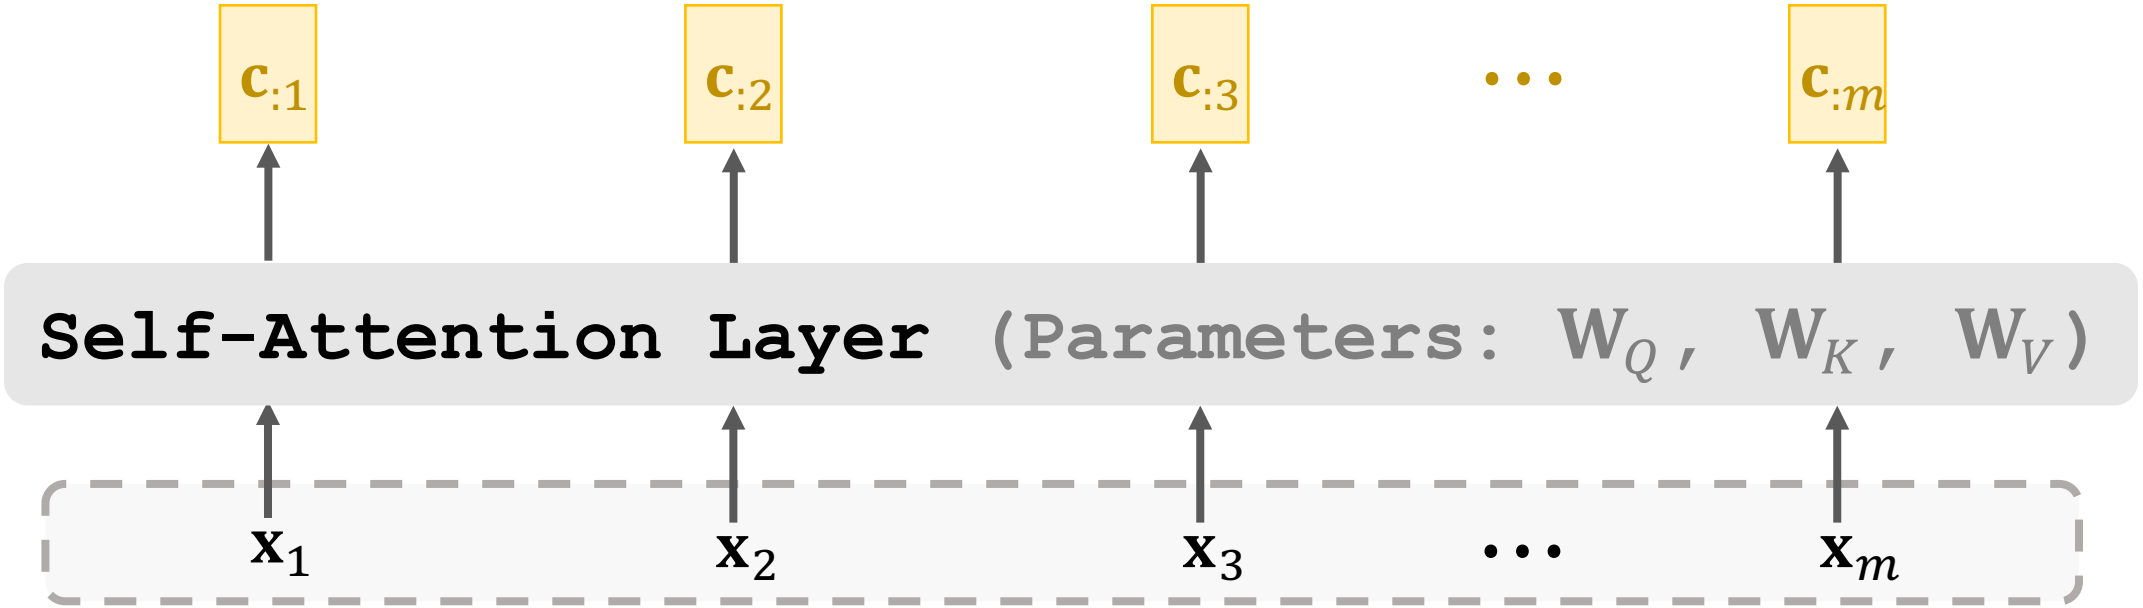

Multi-Head Self-Attention

- Using l single-head self-attentions (which do not share parameters.)
  - A single-head self-attention has 3 parameter matrices: $\mathbf{W}_Q$, $\mathbf{W}_K$, $\mathbf{W}_V$.
  - Totally $3l$ parameters matrices.

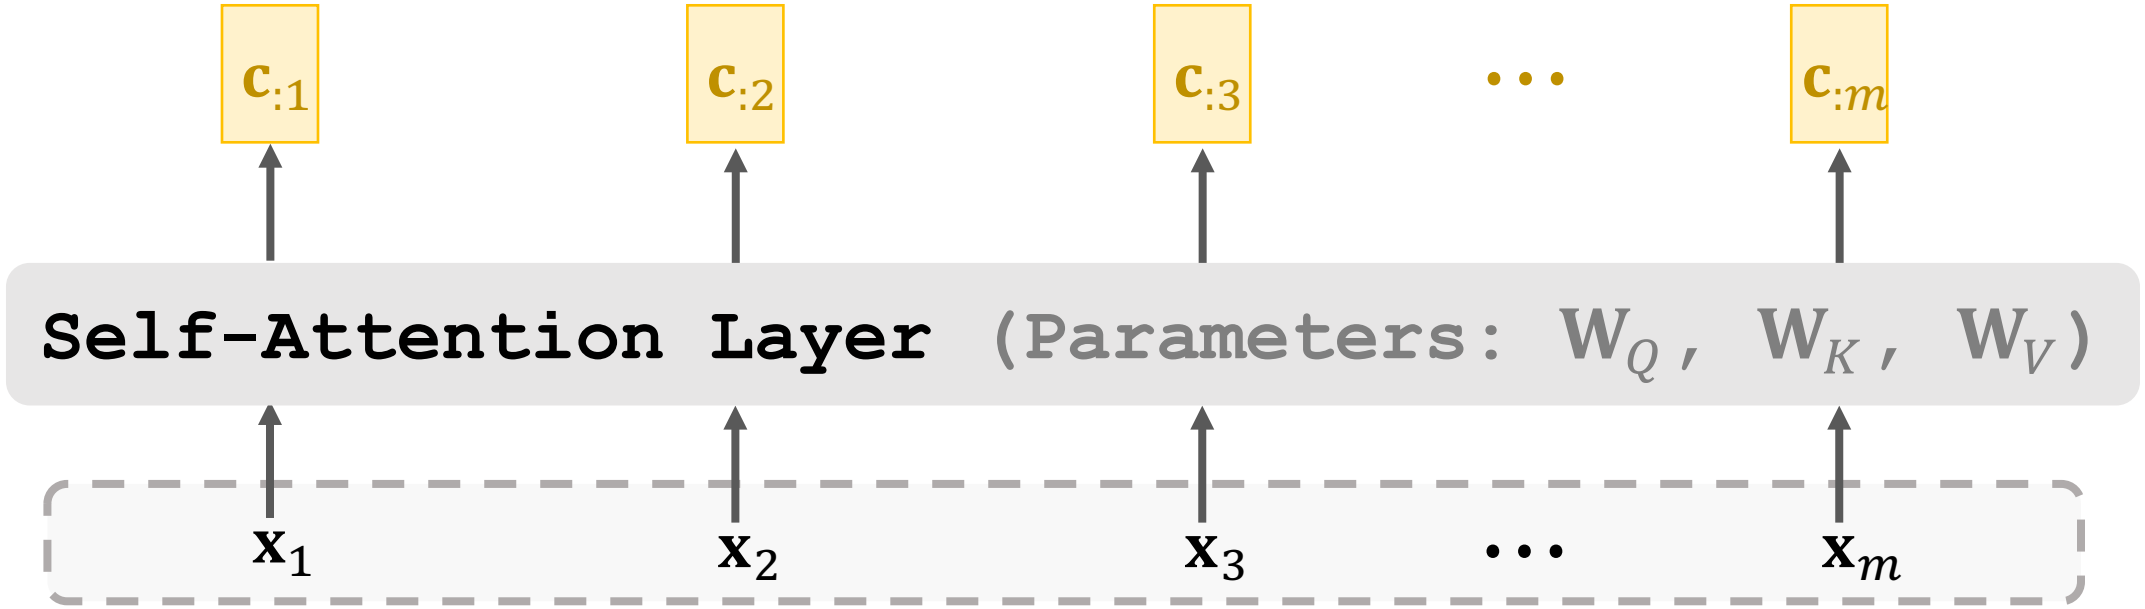

Multi-Head Self-Attention

- Using l single-head self-attentions (which do not share parameters.)
- Concatenating outputs of single-head self-attentions.
  - Suppose single-head self-attentions' outputs are $d \times m$ matrices.
  - Multi-head's output shape: $(ld) \times m$.

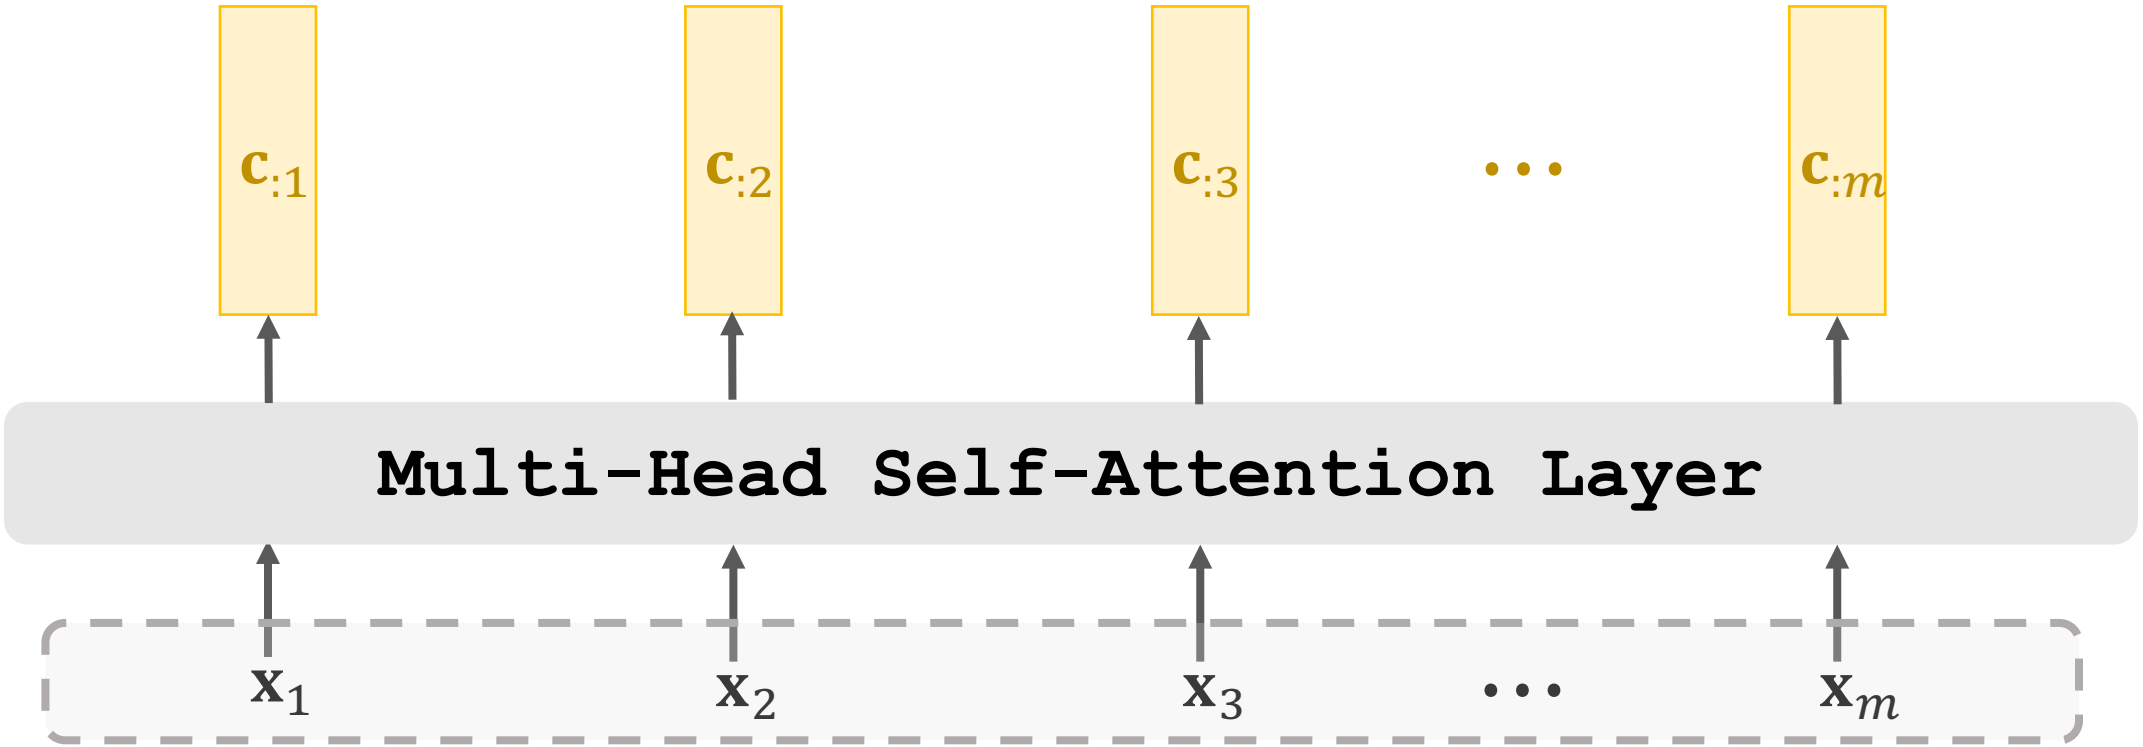

Multi-Head Attention

- Using l single-head attentions (which do not share parameters.)
- Concatenating single-head attentions' outputs.

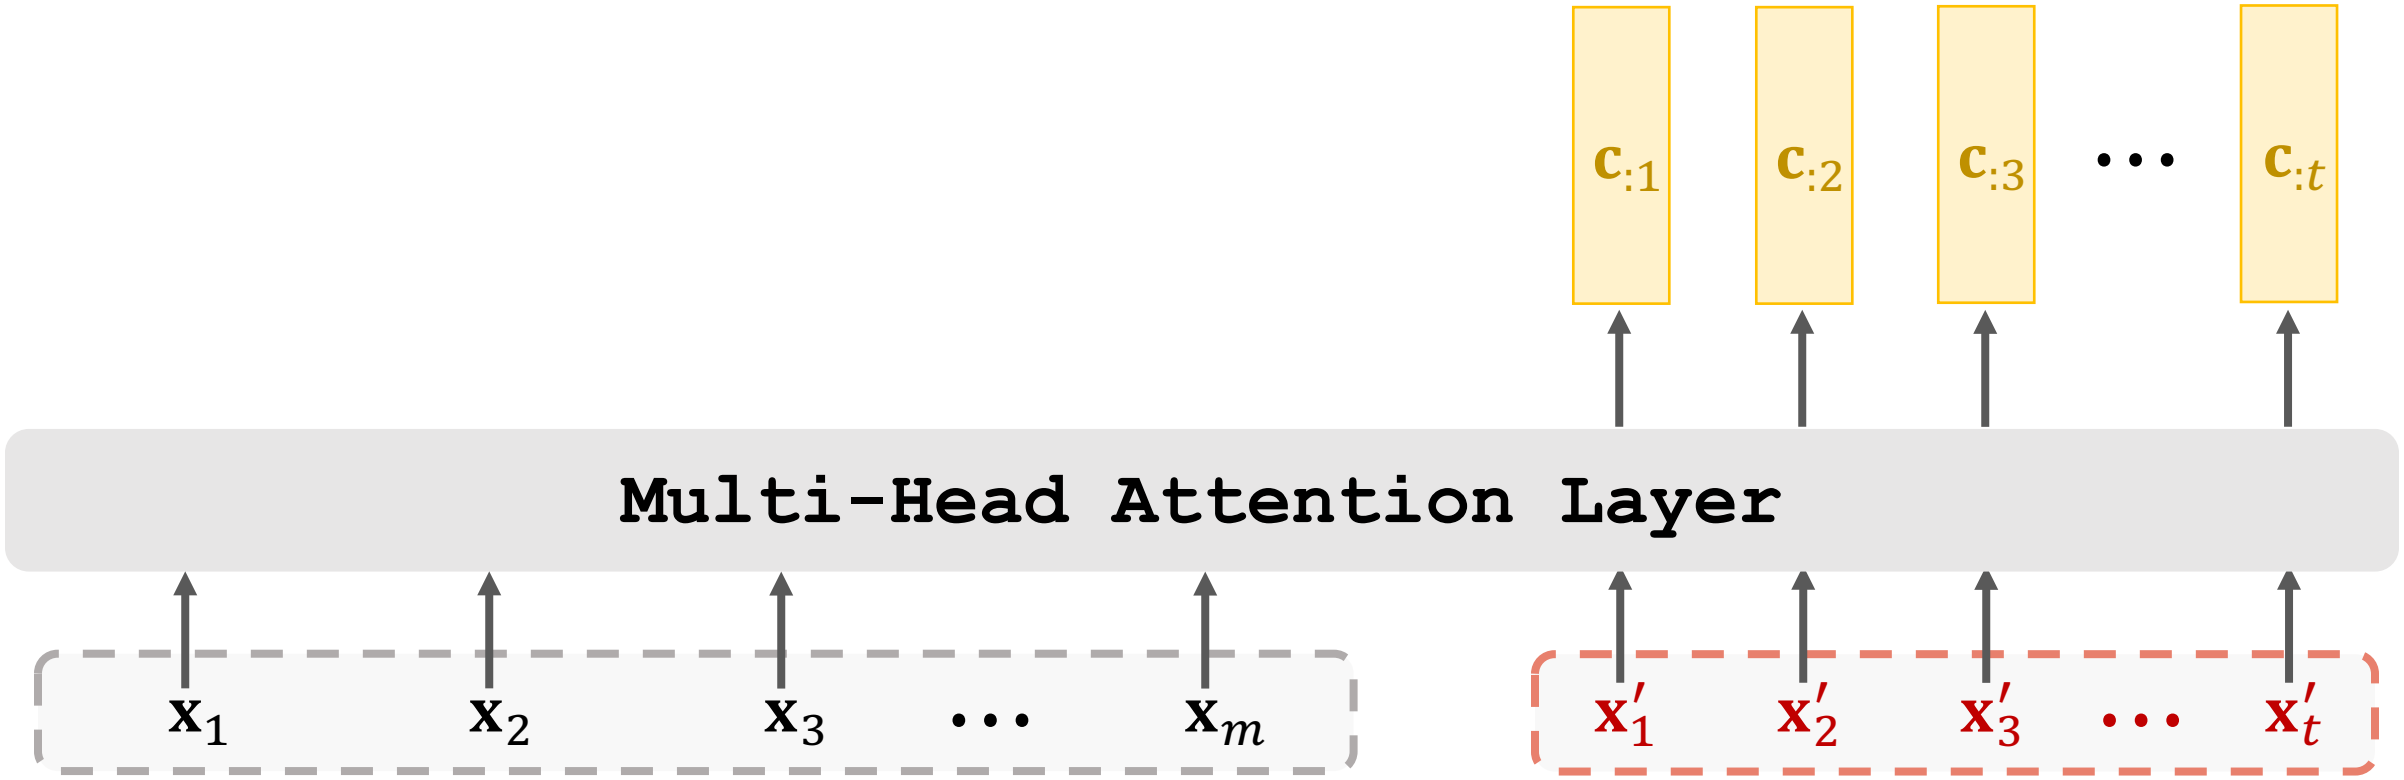

Stacked Self-Attention Layers

Self-Attention Layer

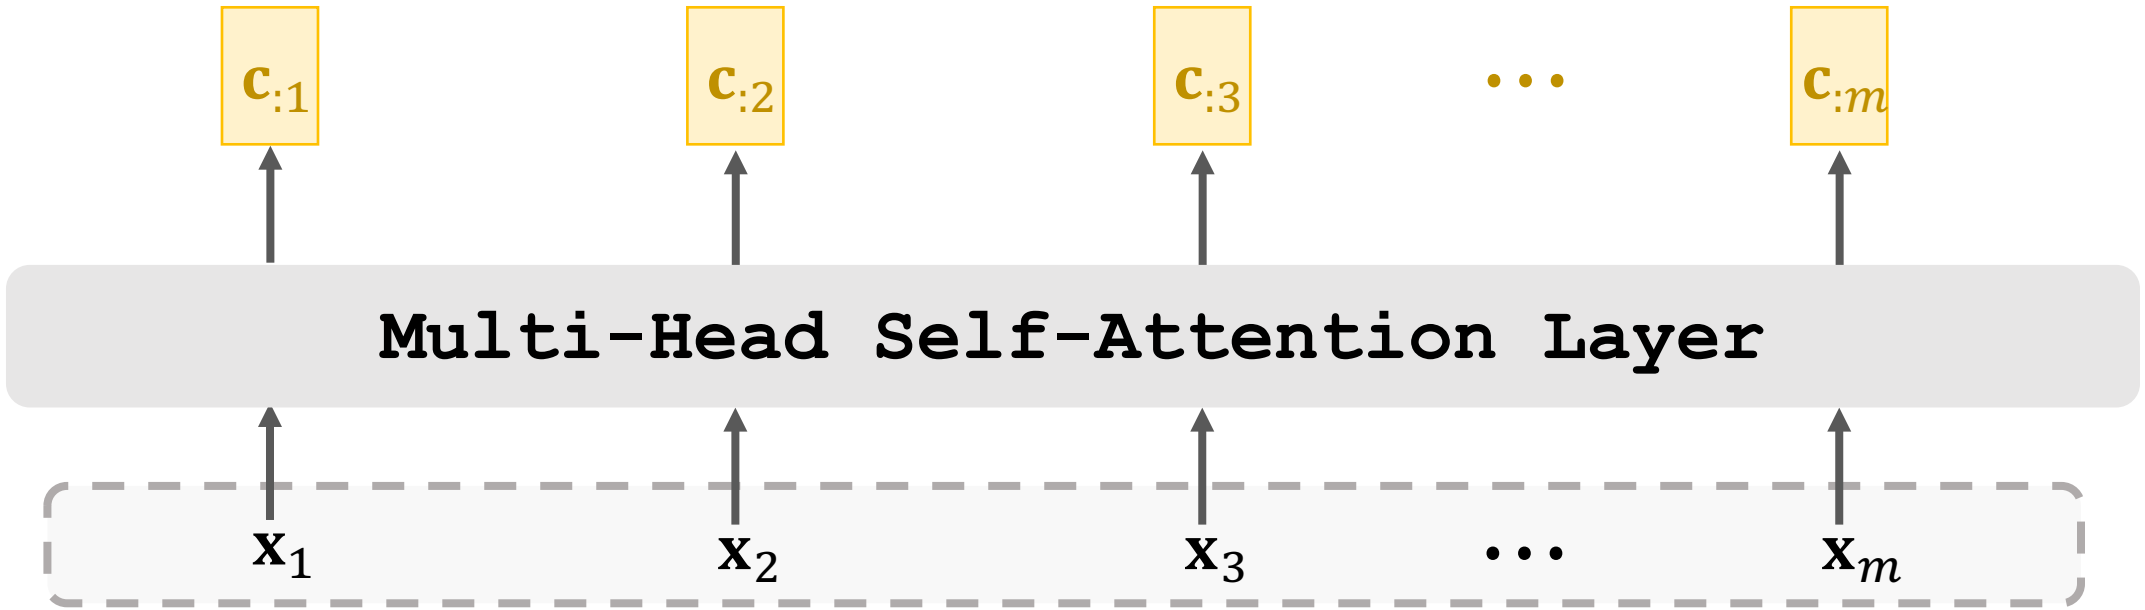

Self-Attention Layer + Dense Layer

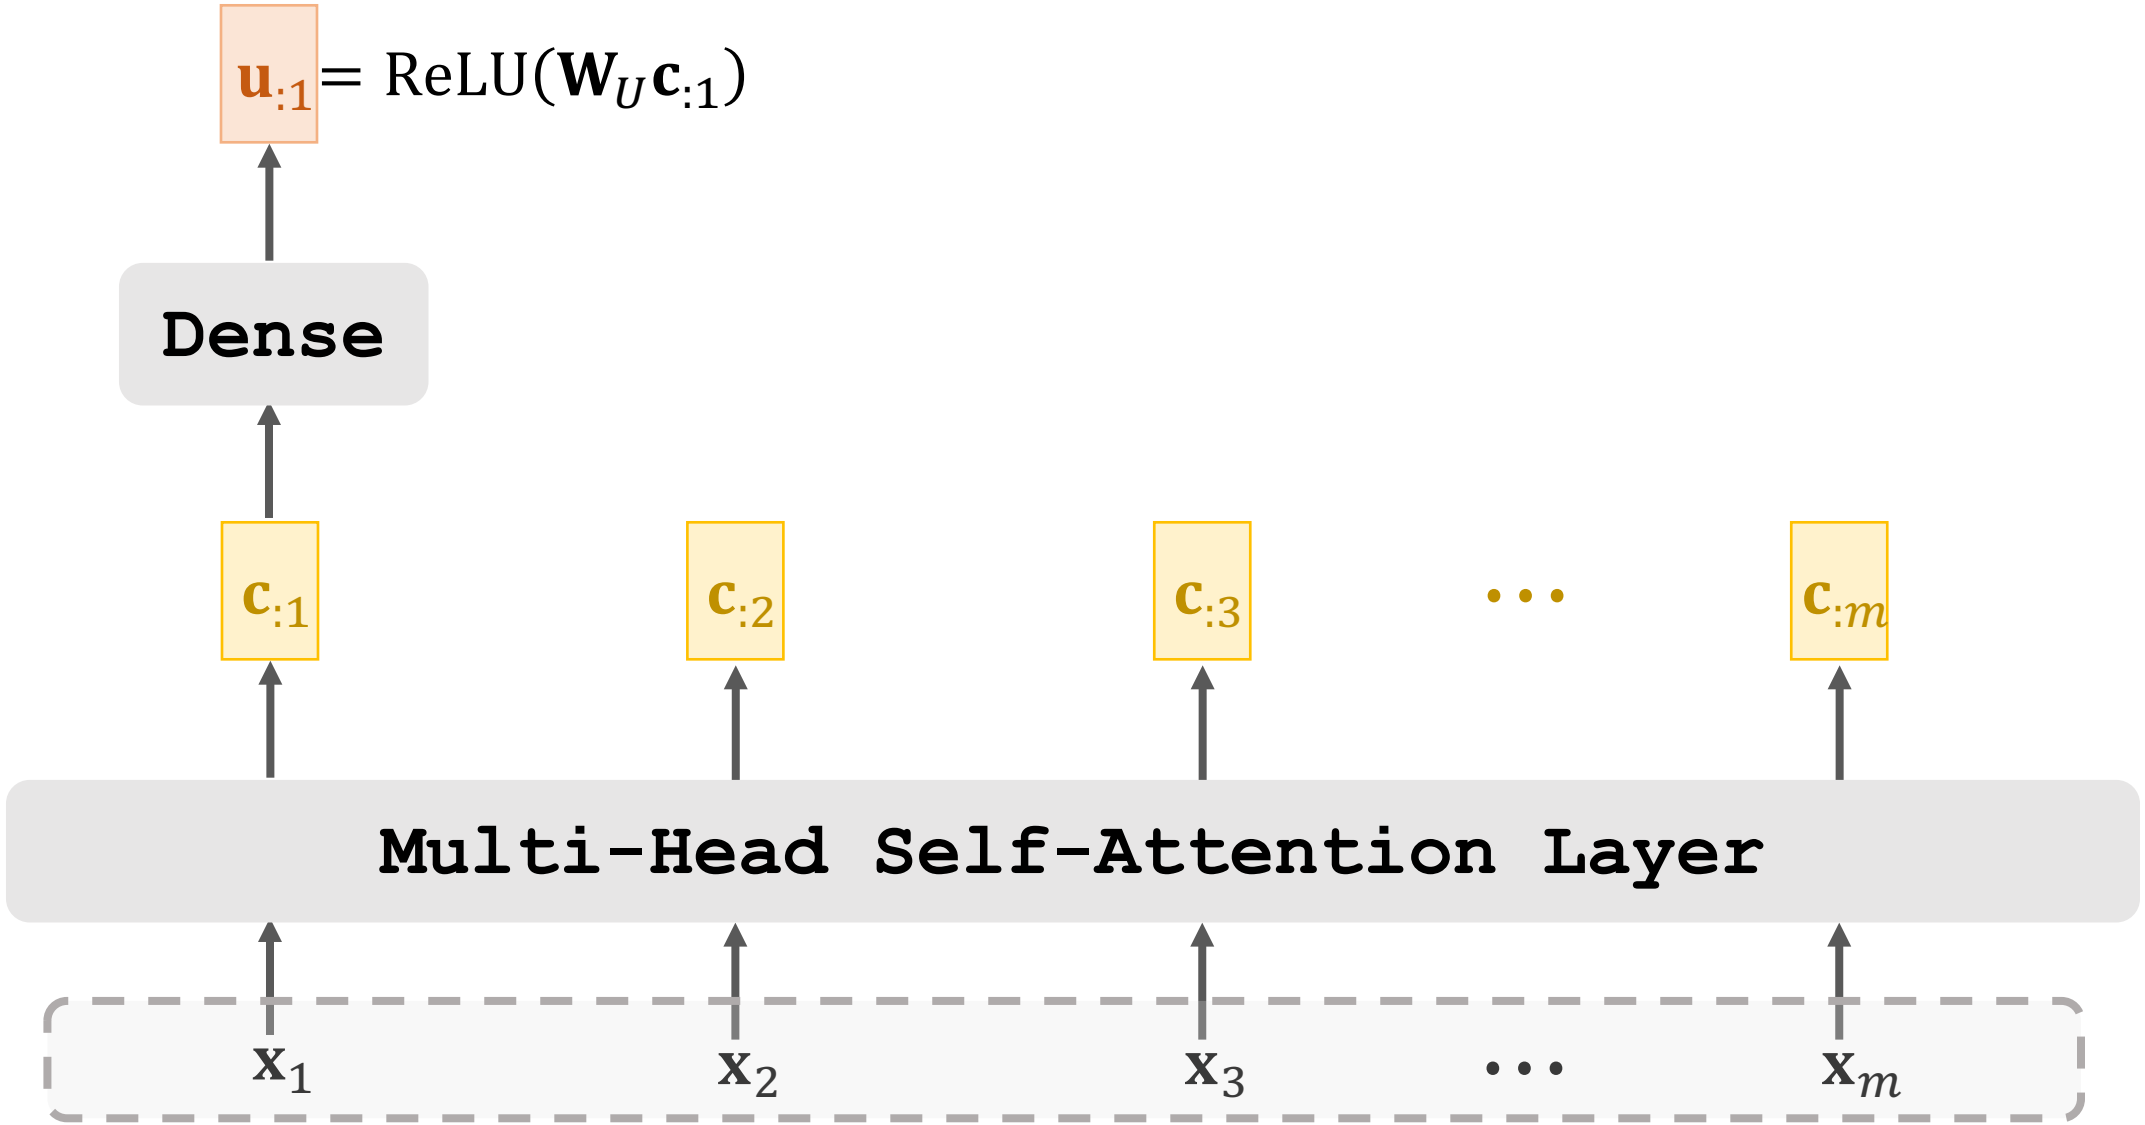

Self-Attention Layer + Dense Layer

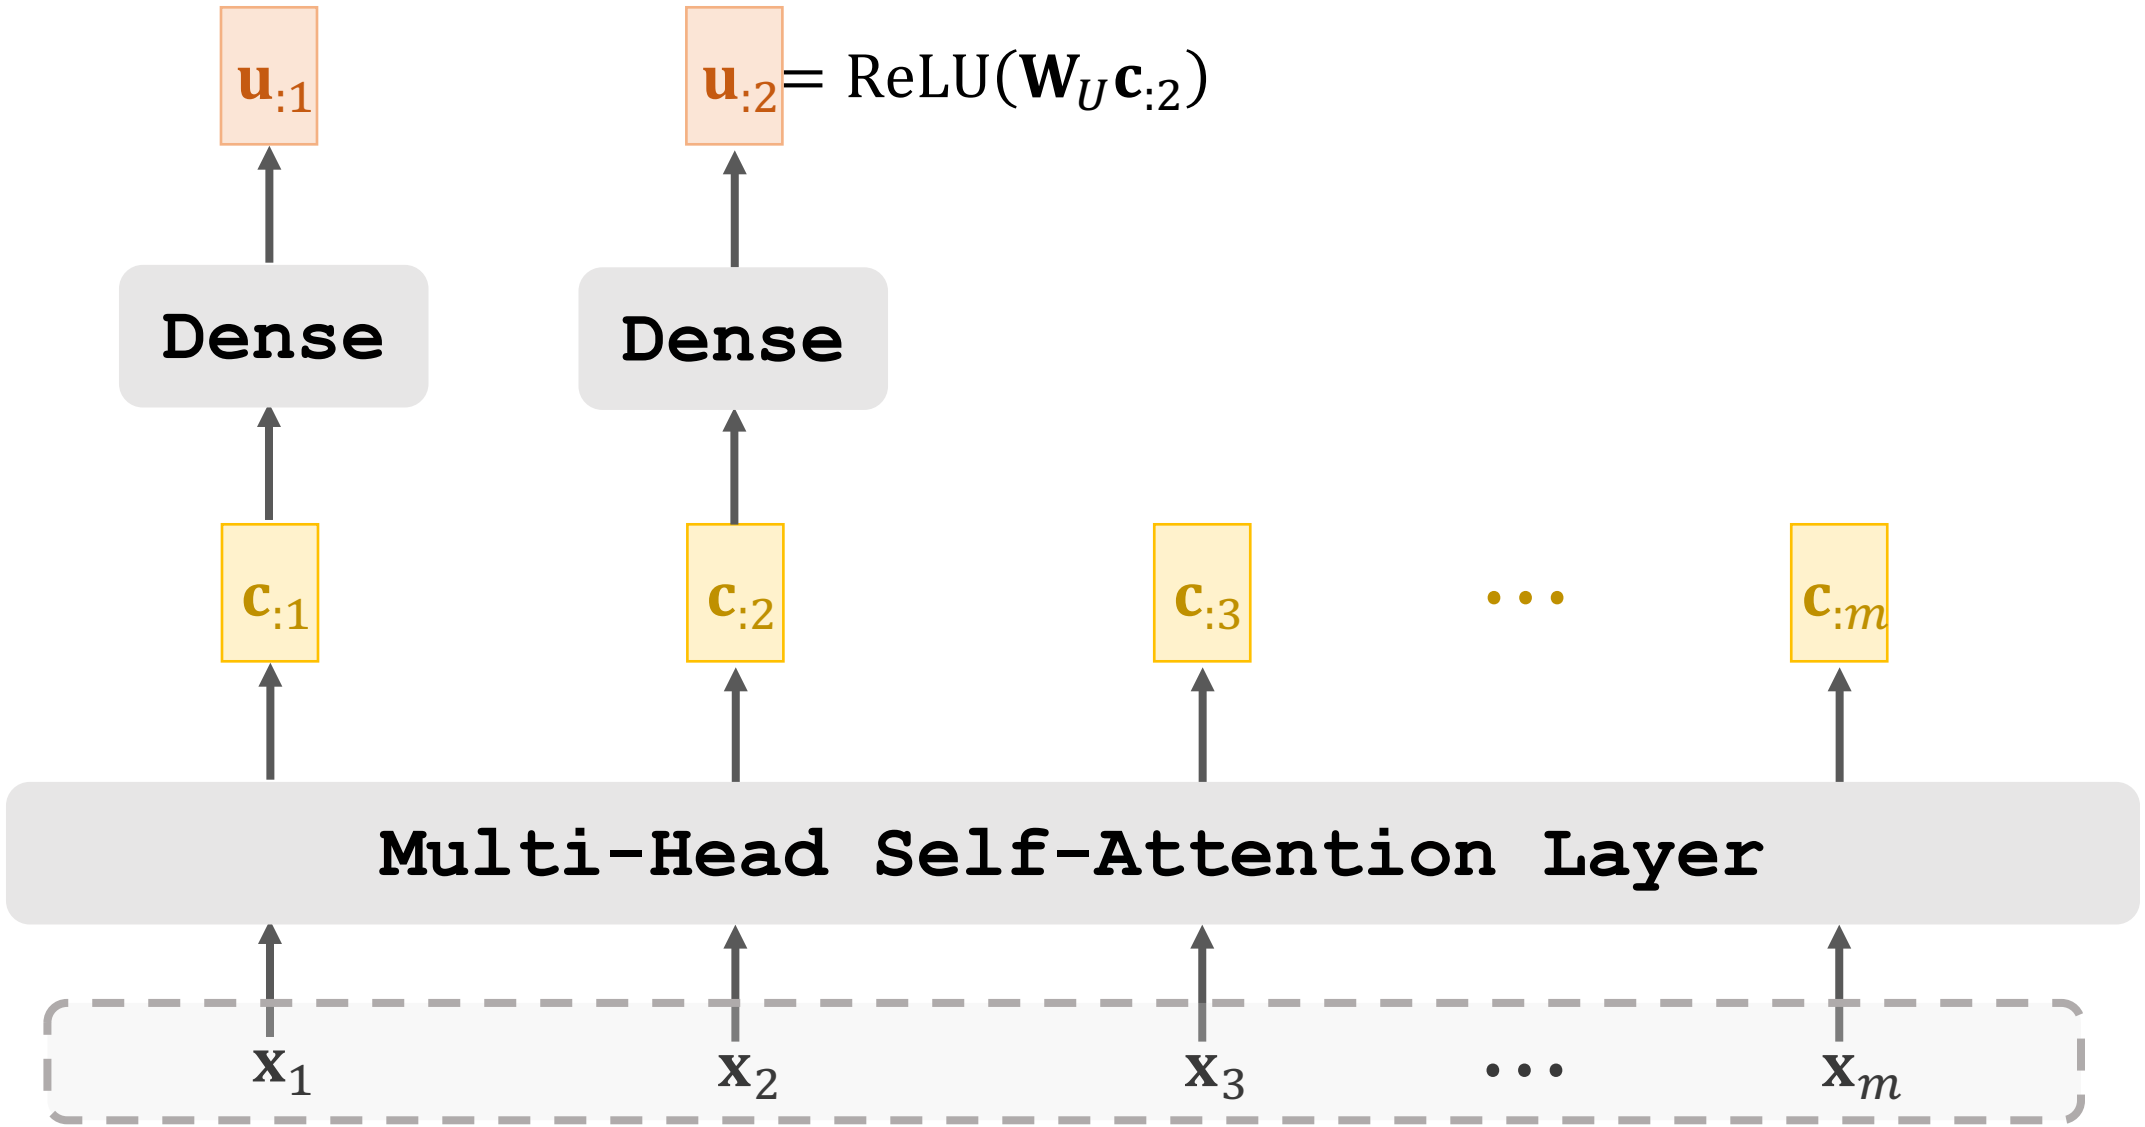

Self-Attention Layer + Dense Layer

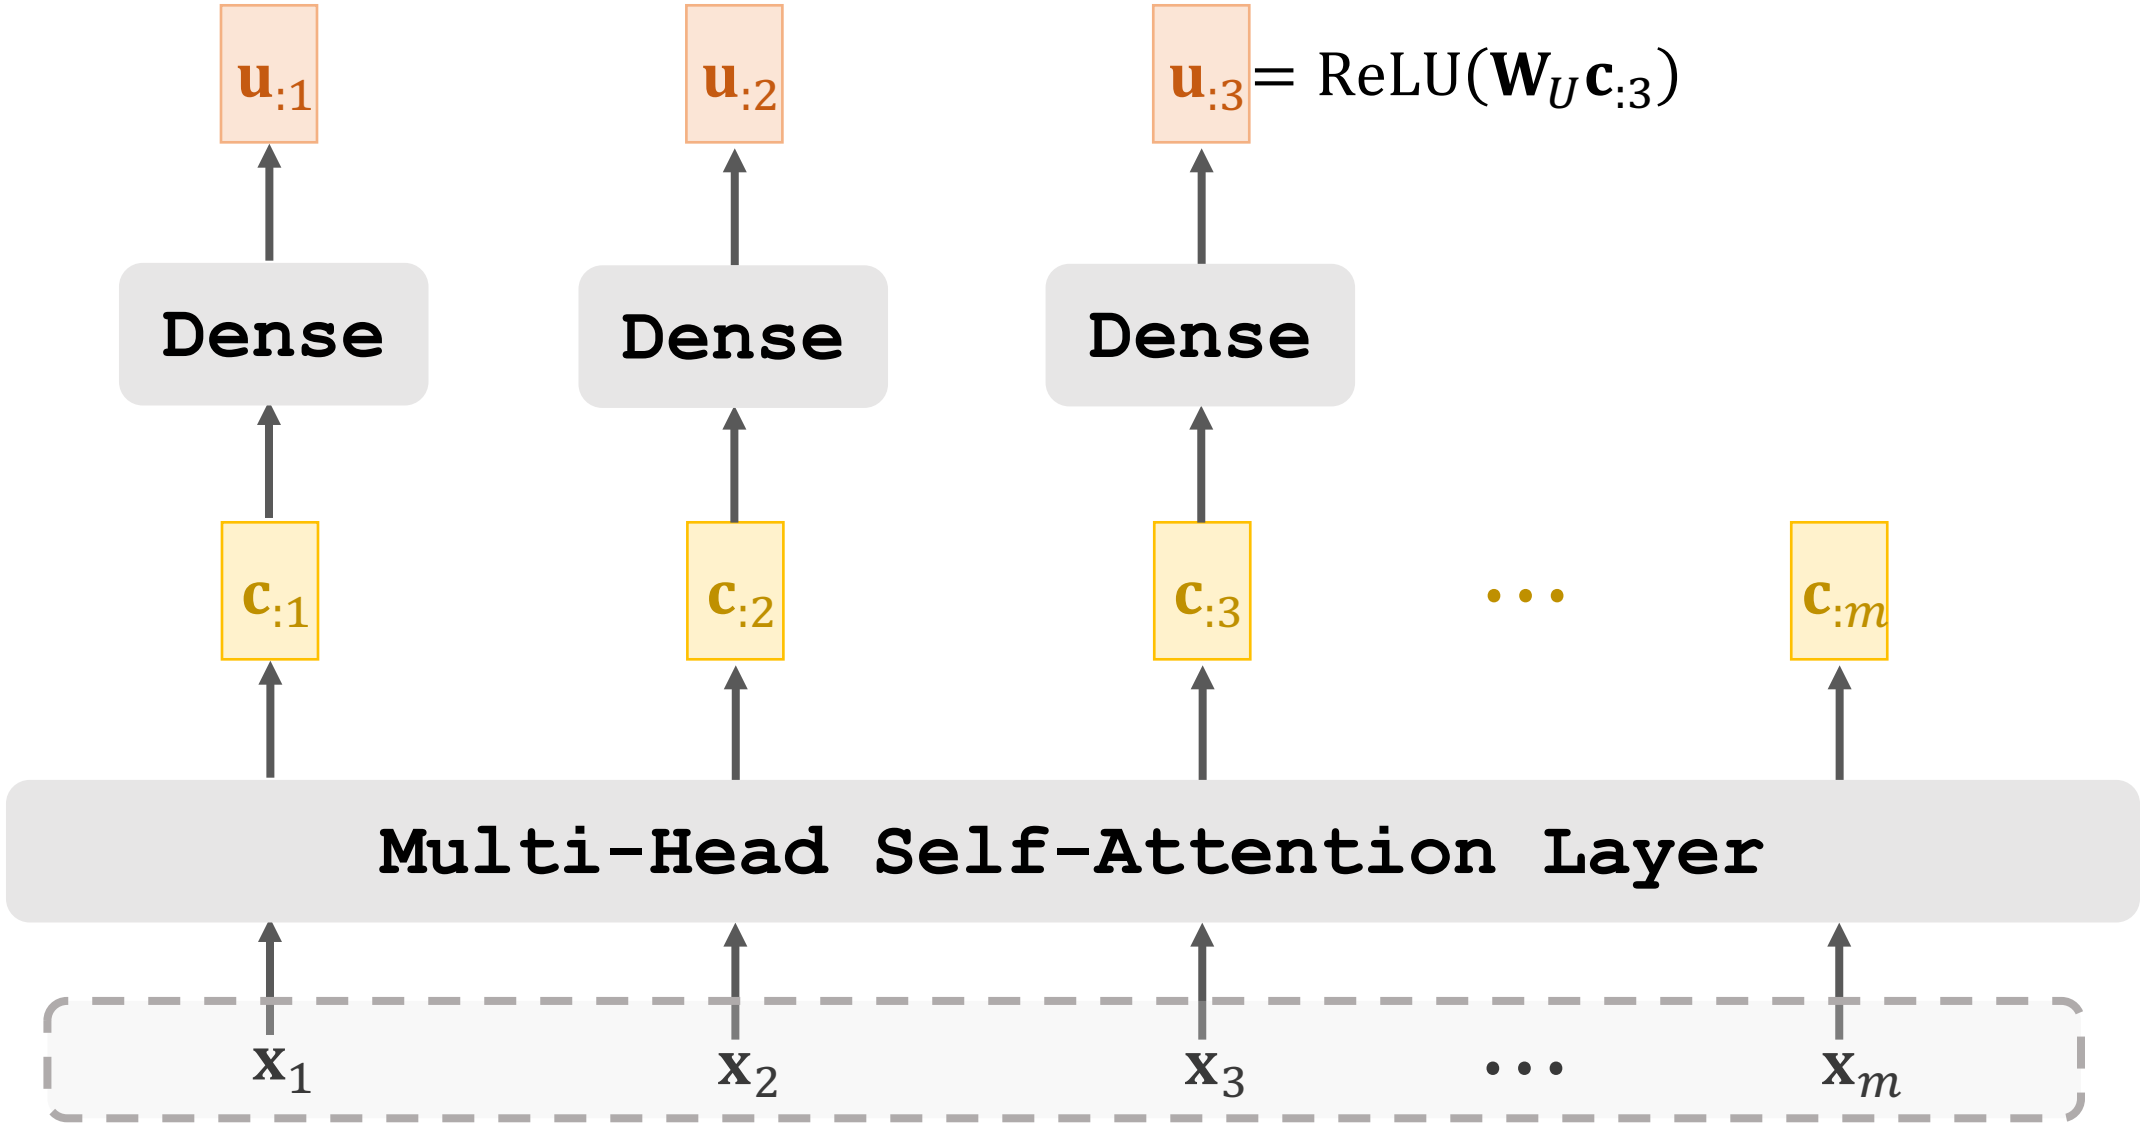

Self-Attention Layer + Dense Layer

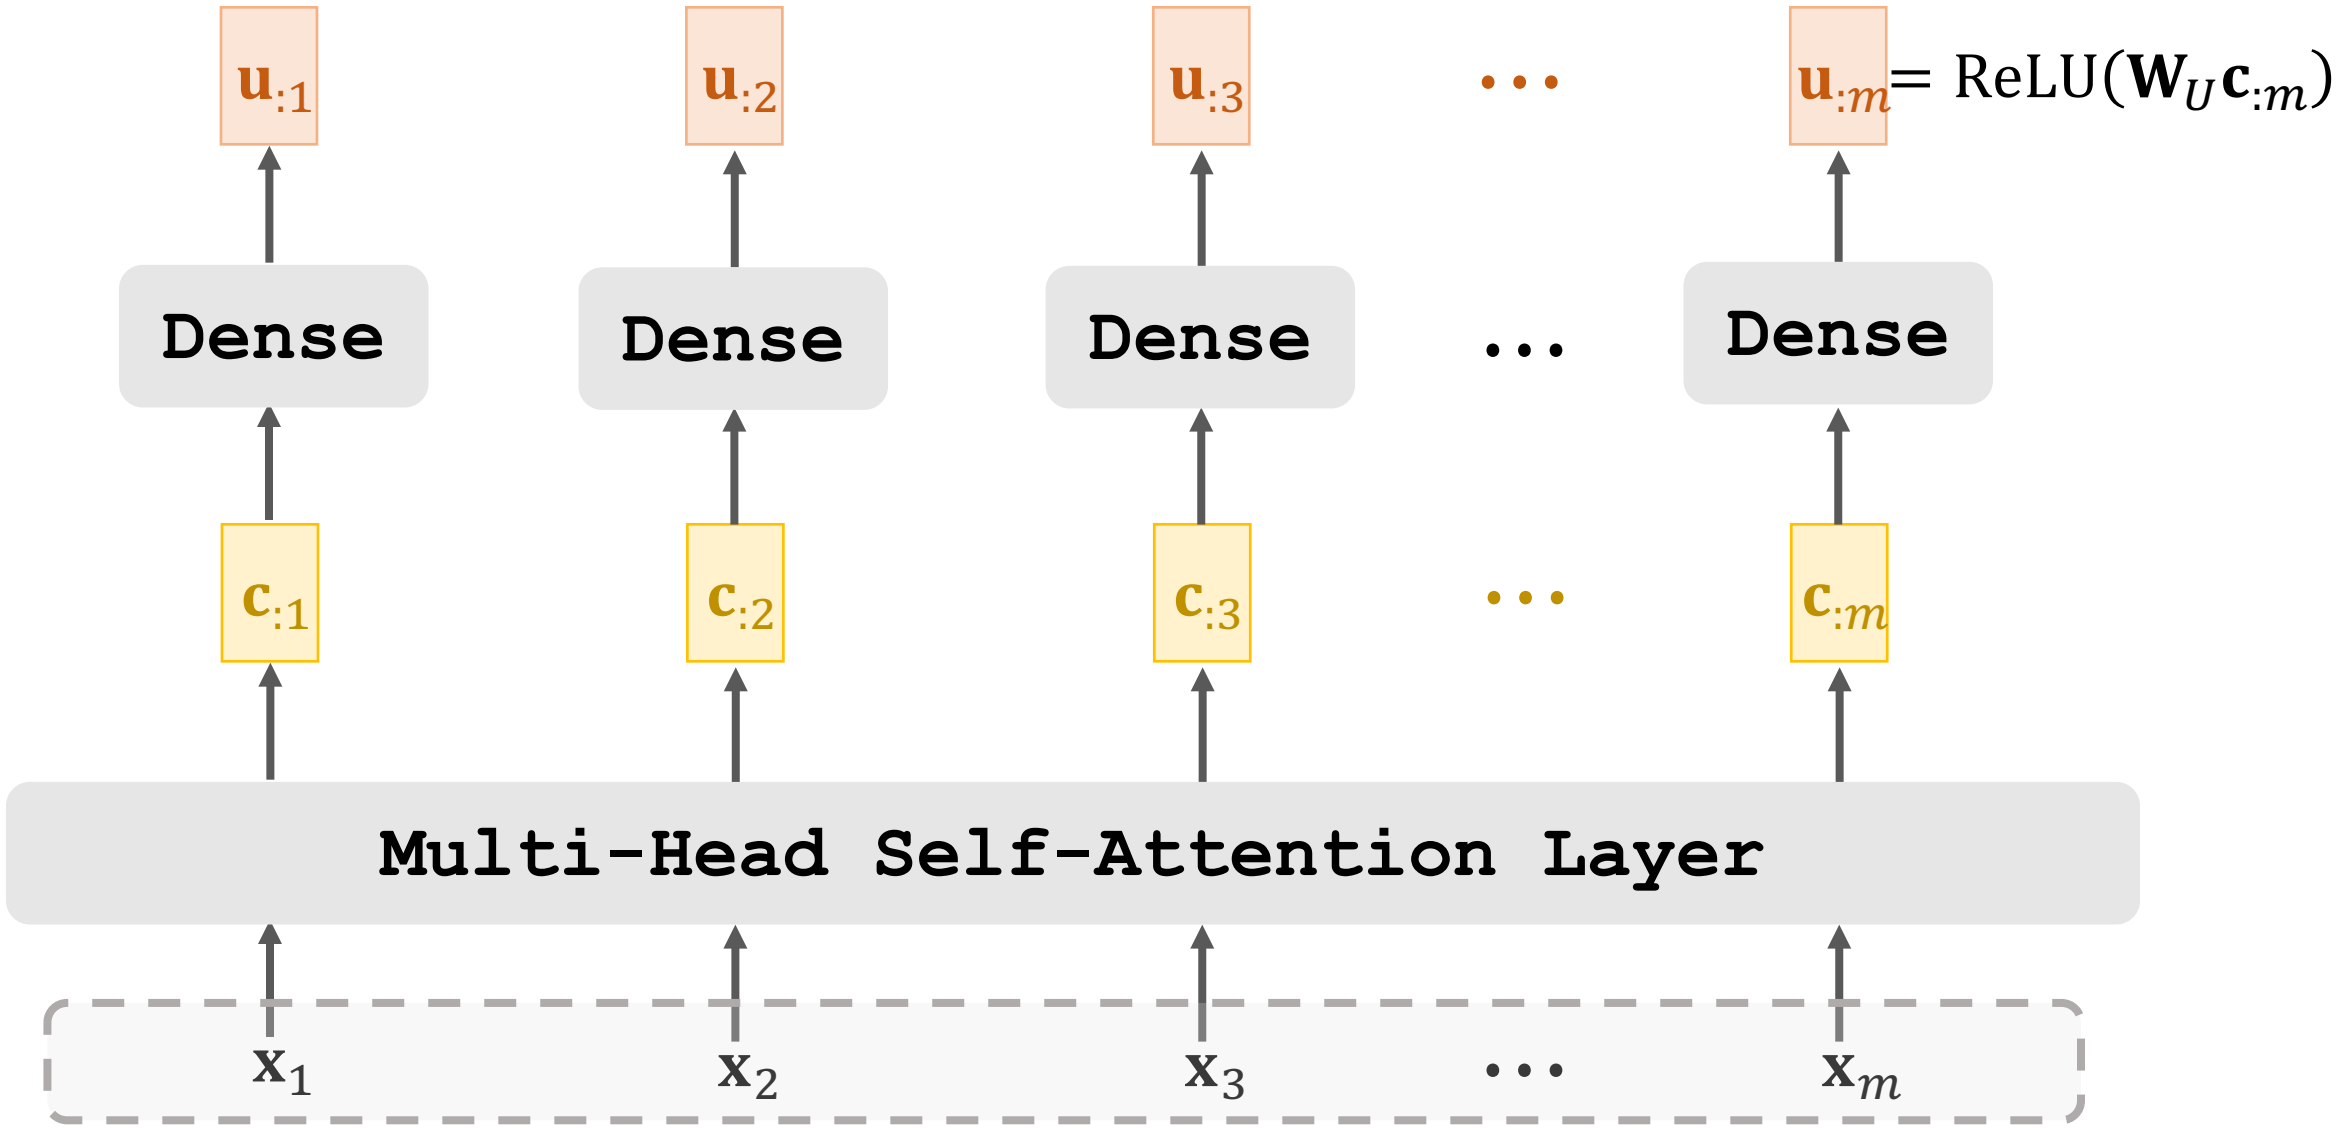

Stacked Self-Attention Layers

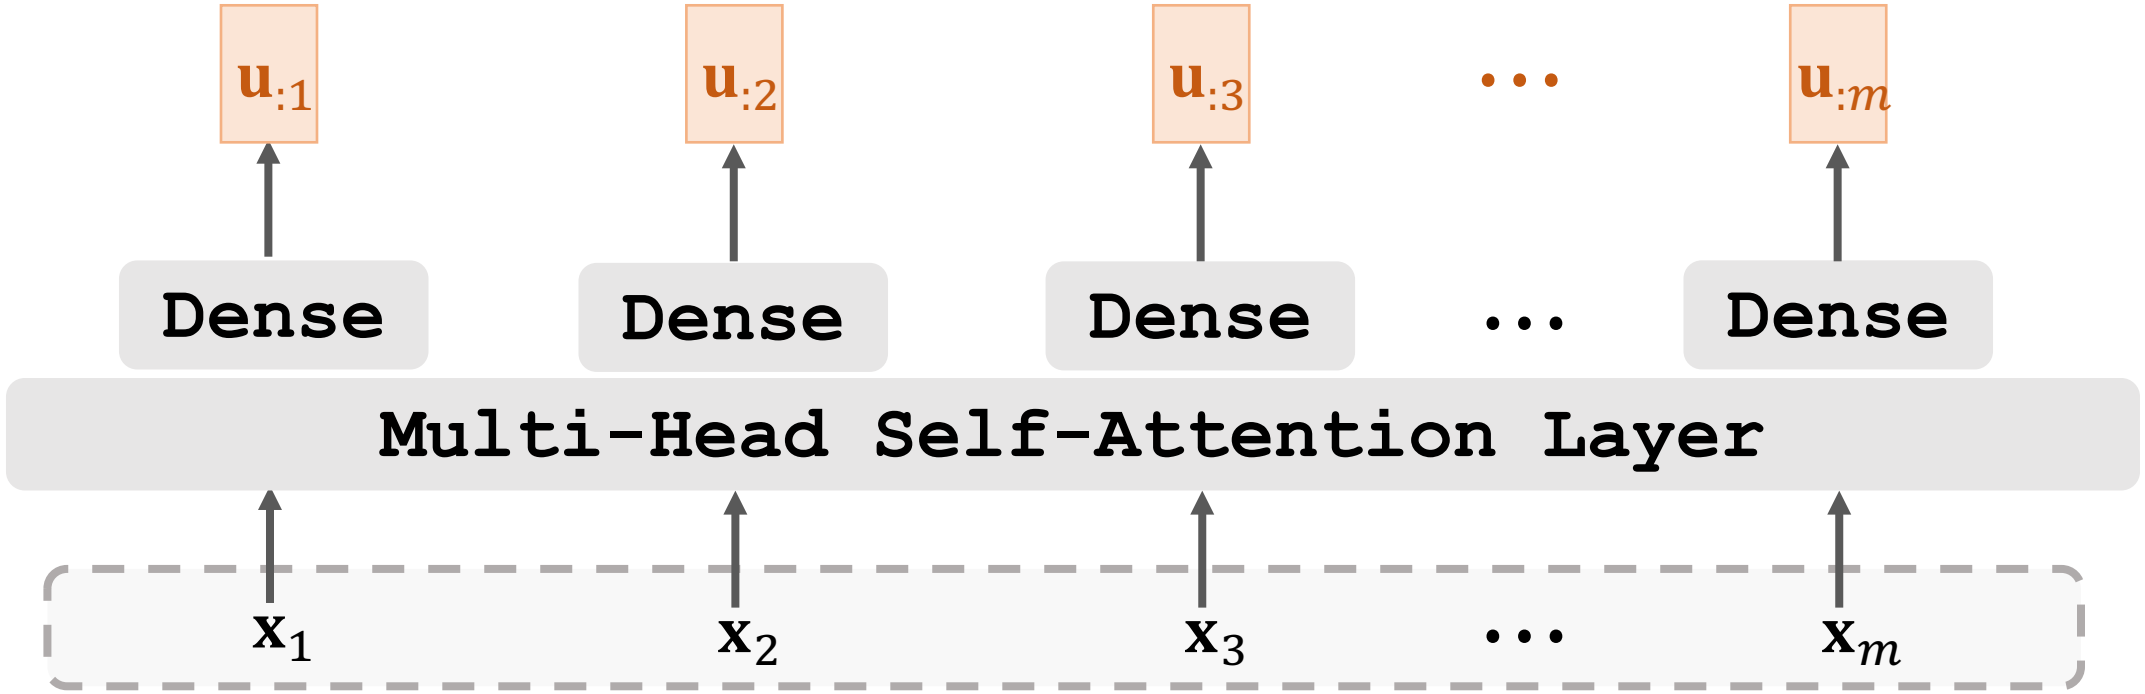

Stacked Self-Attention Layers

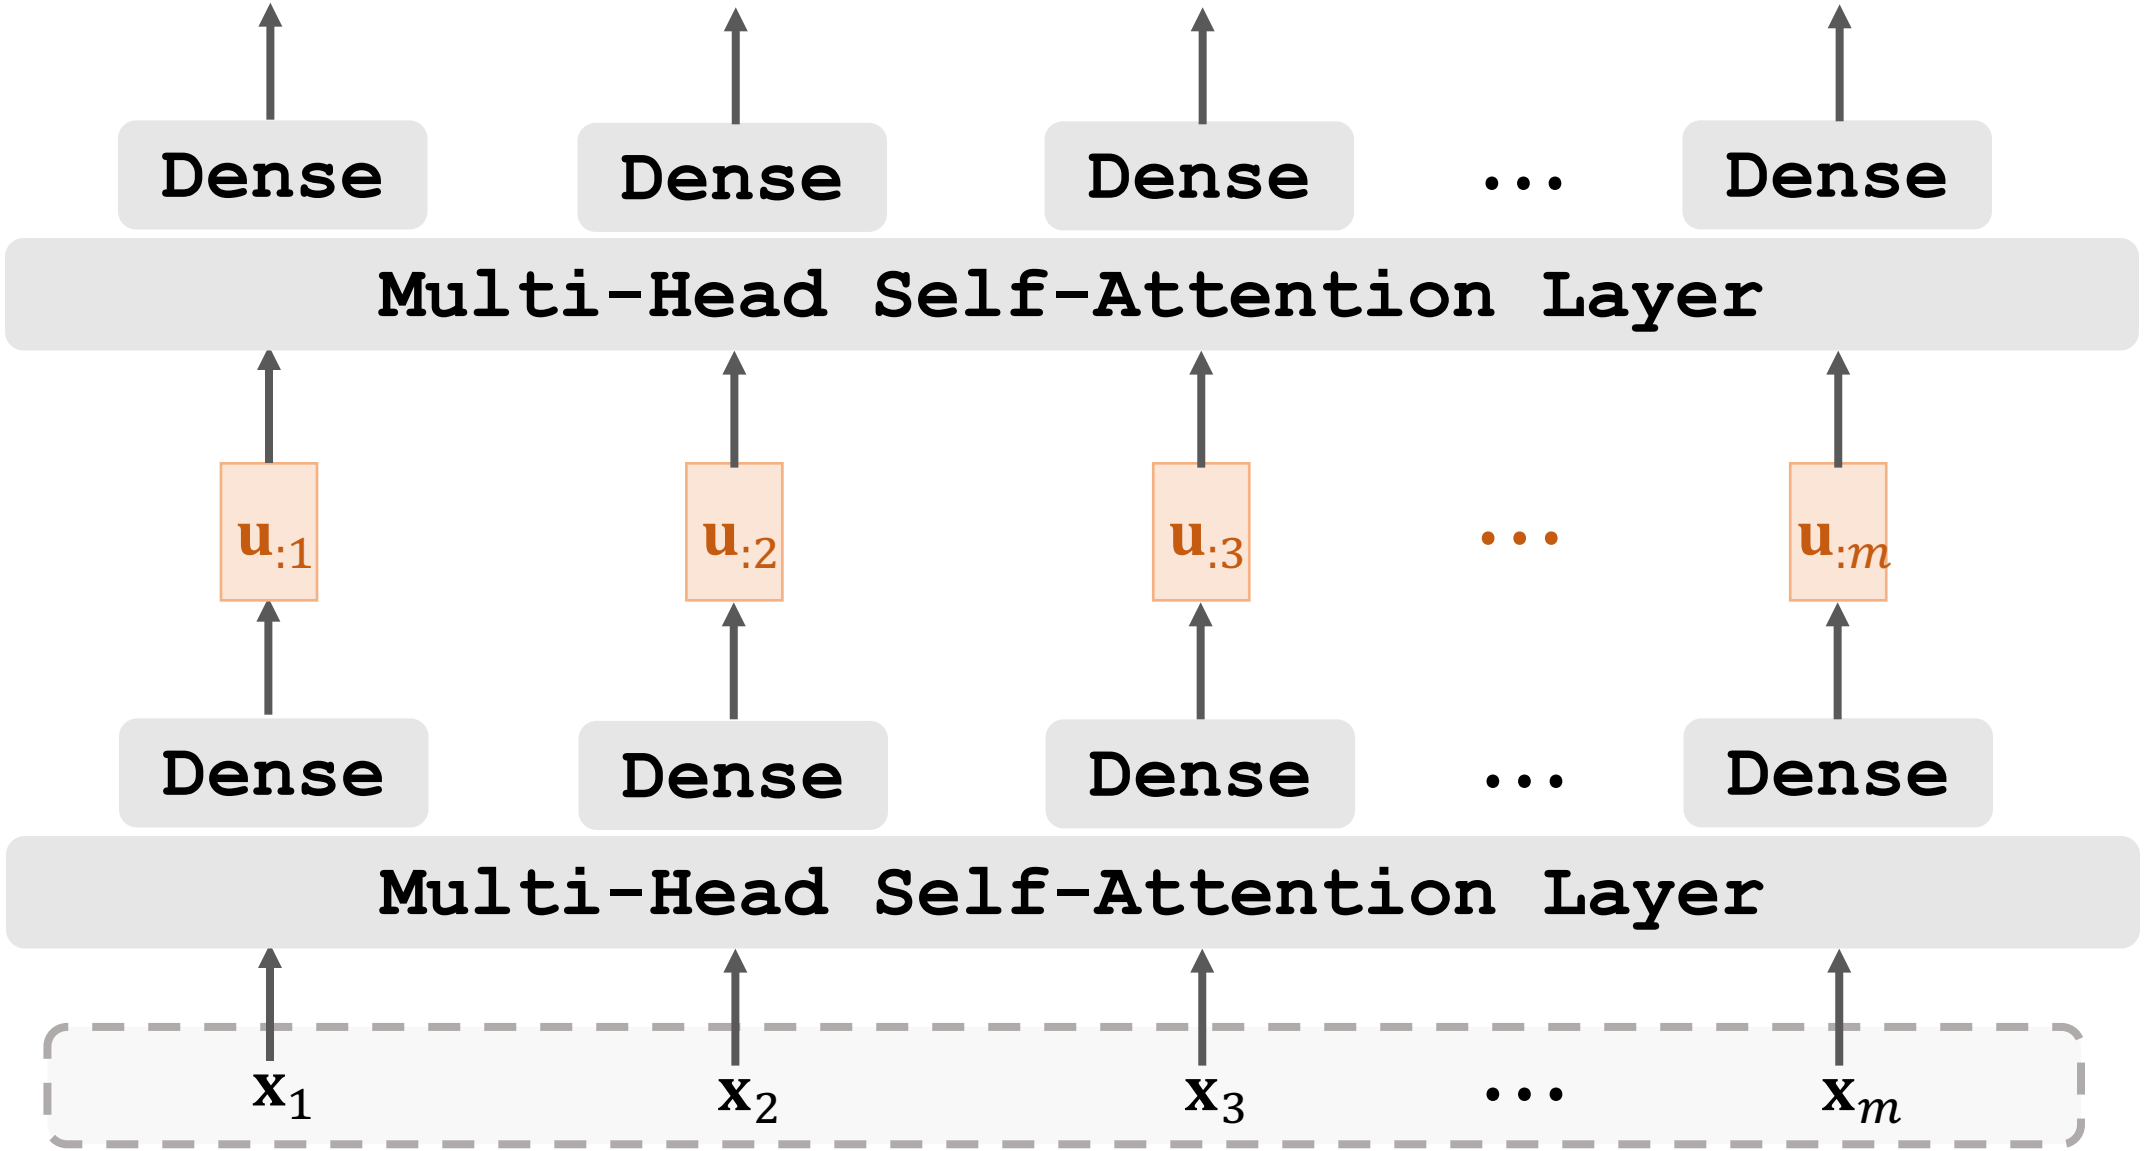

Transformer's Encoder

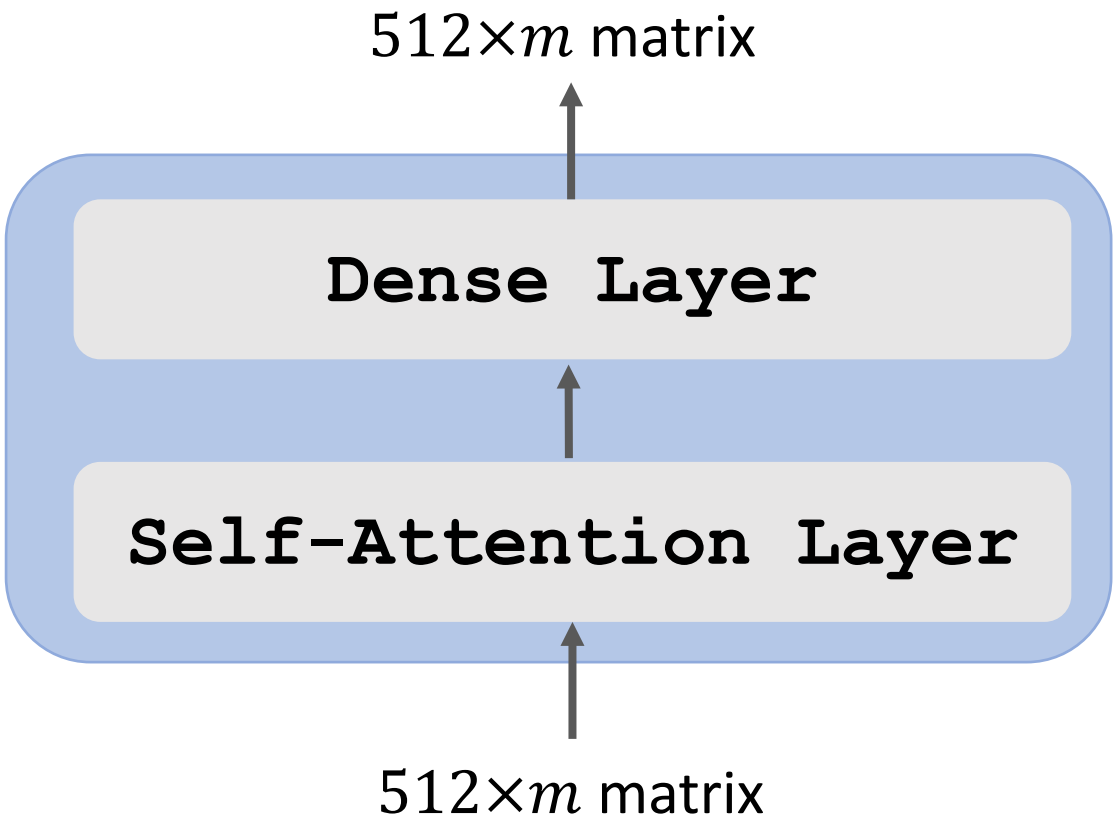

One Block of Encoder

Transformer's Encoder

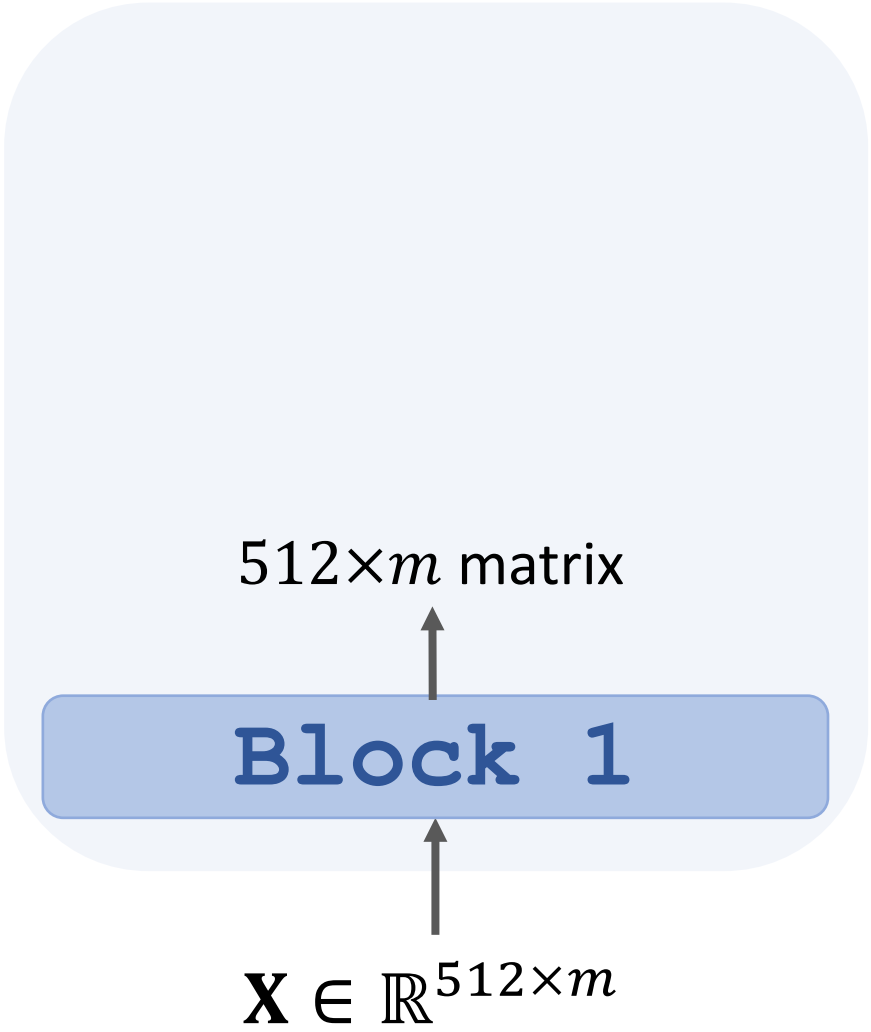

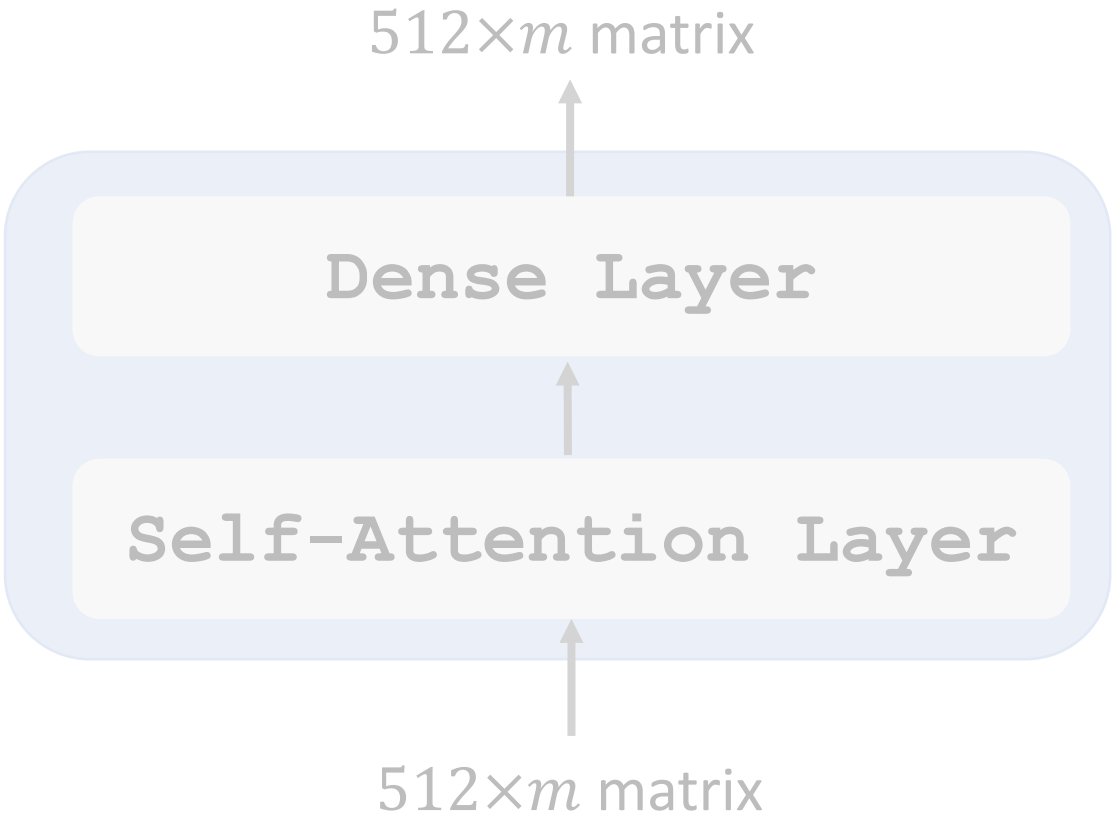

One Block of Encoder

Transformer's Encoder

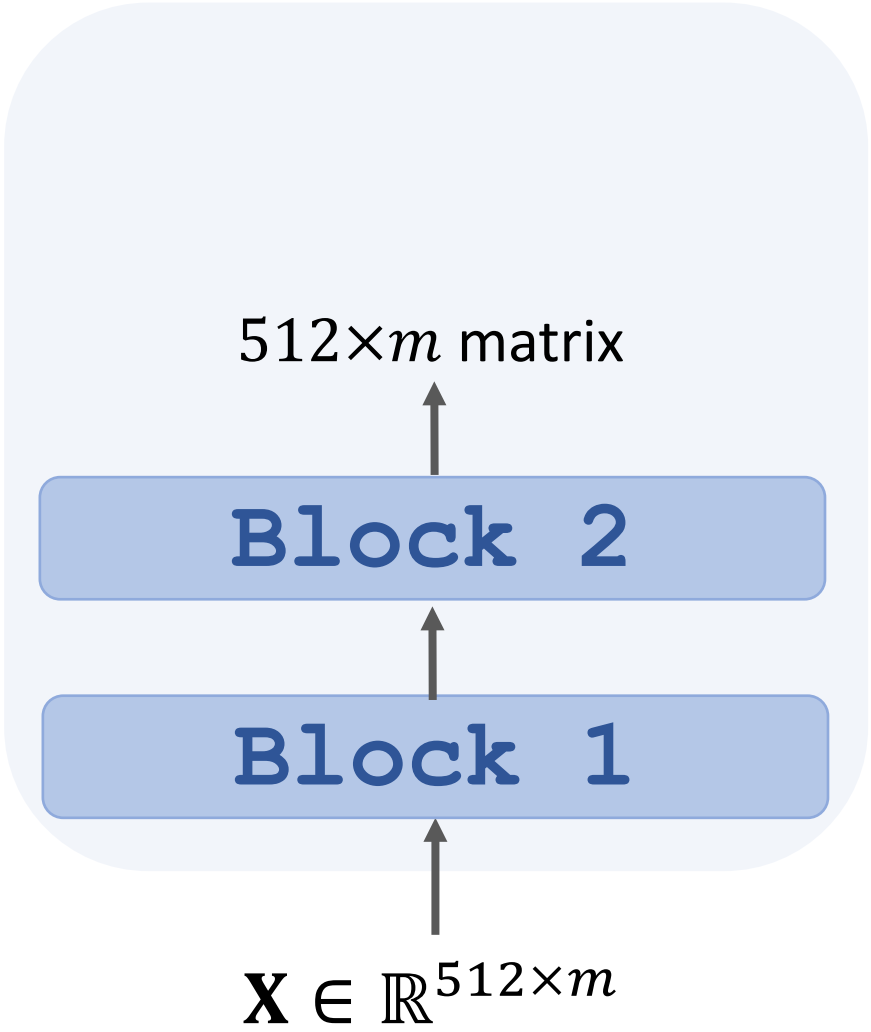

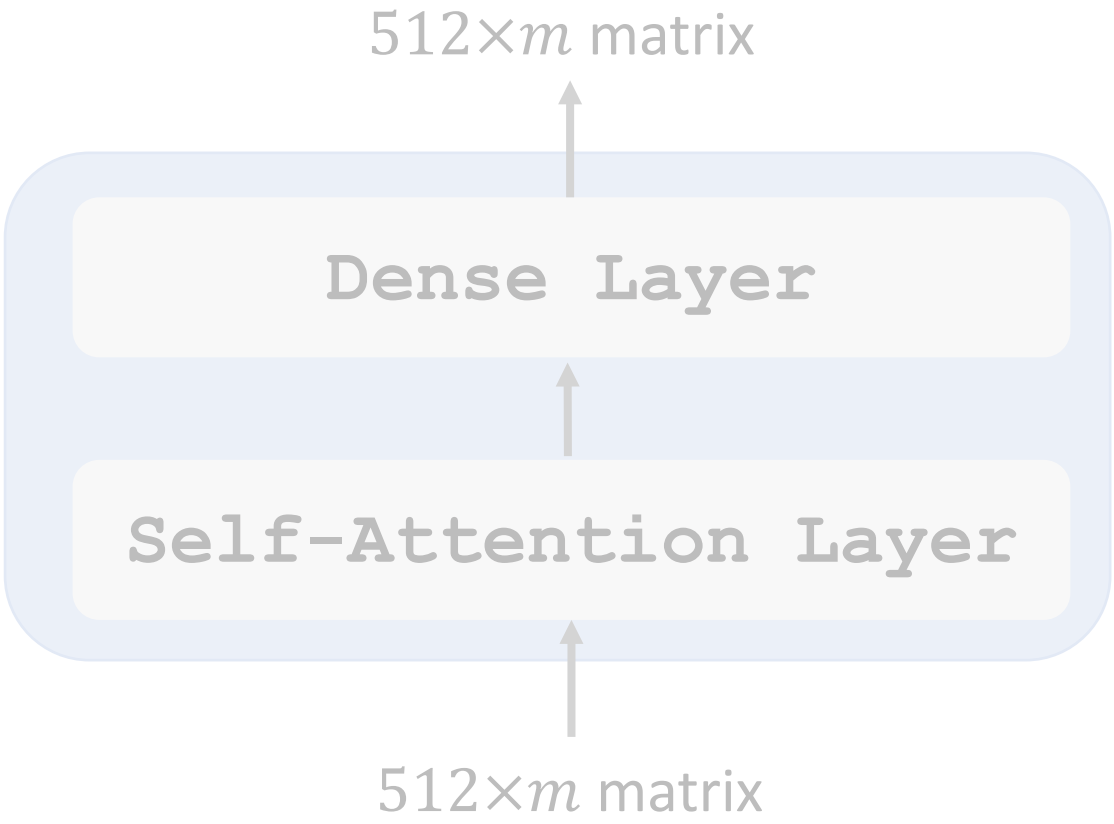

One Block of Encoder

Transformer's Encoder

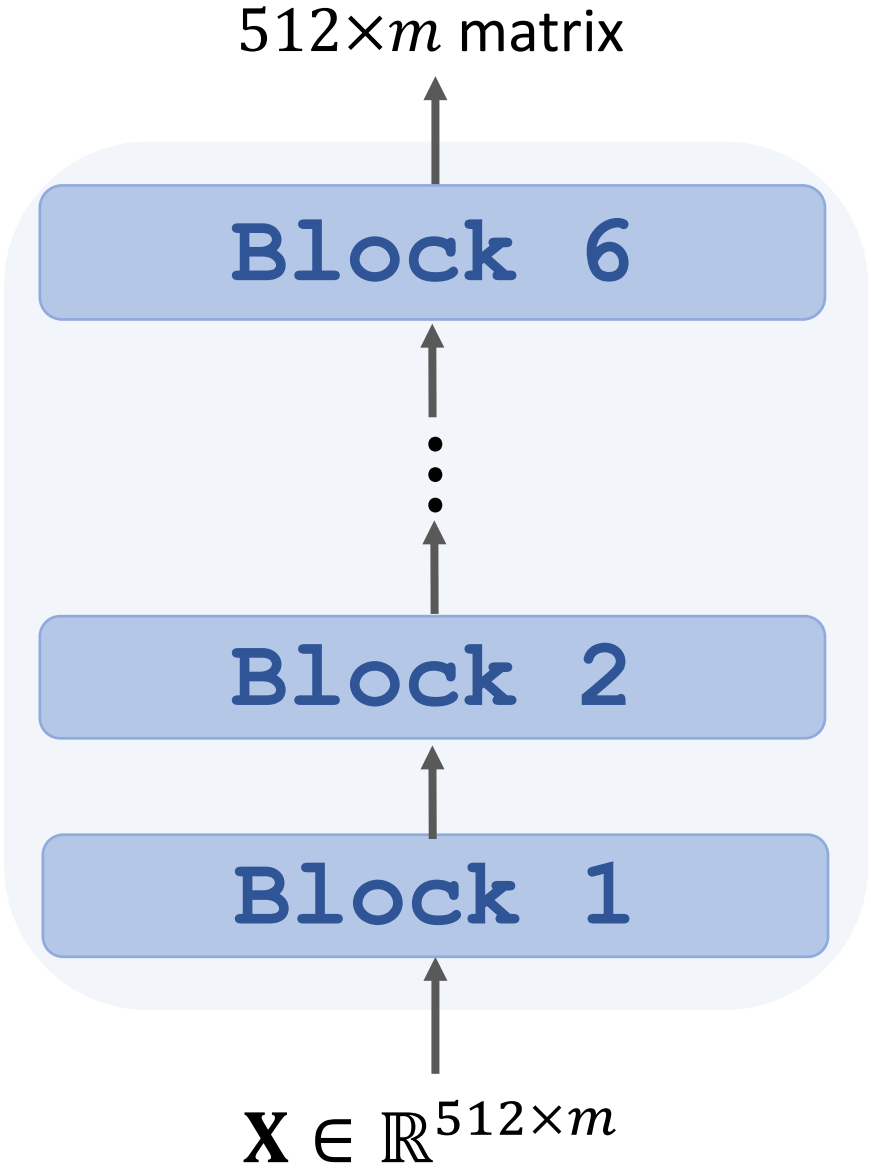

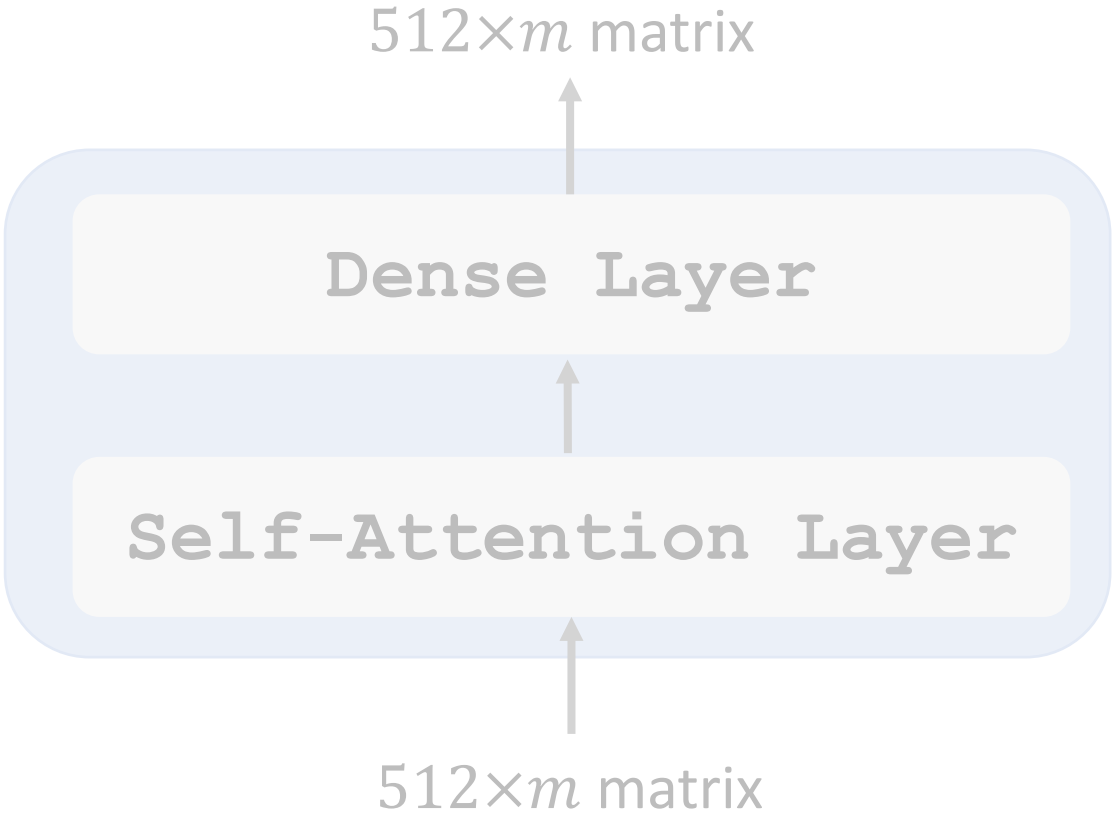

One Block of Encoder

Stacked Attention Layers

Stacked Attentions

- Transformer is a Seq2Seq model (encoder + decoder).
- Encoder's inputs are vectors $\mathbf{x}_1, \mathbf{x}_2, \cdots, \mathbf{x}_m$.
- Decoder's inputs are vectors $\mathbf{x}'_1, \mathbf{x}'_2, \cdots, \mathbf{x}'_t$.

Encoder's inputs:

$\mathbf{x}_1$ $\mathbf{x}_2$ $\mathbf{x}_3$ $\cdots$ $\mathbf{x}_m$

Decoder's inputs:

$\mathbf{x}'_1$ $\mathbf{x}'_2$ $\mathbf{x}'_3$ $\cdots$ $\mathbf{x}'_t$

Stacked Attentions

- Transformer's encoder contains 6 stacked blocks.
- 1 block ≈ 1 multi-head attention layer + 1 dense layer.

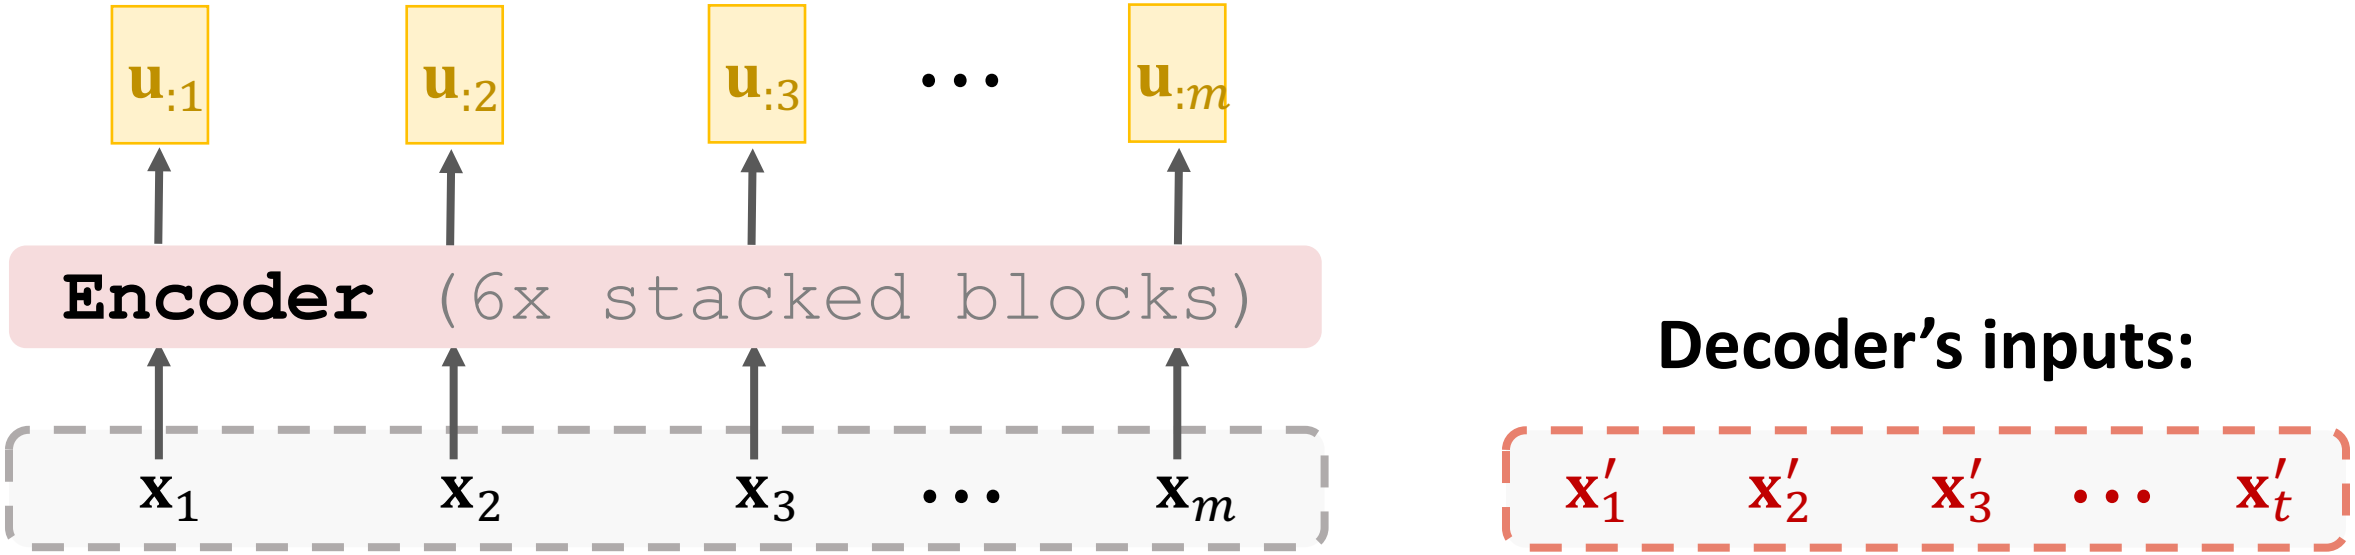

Stacked Attentions

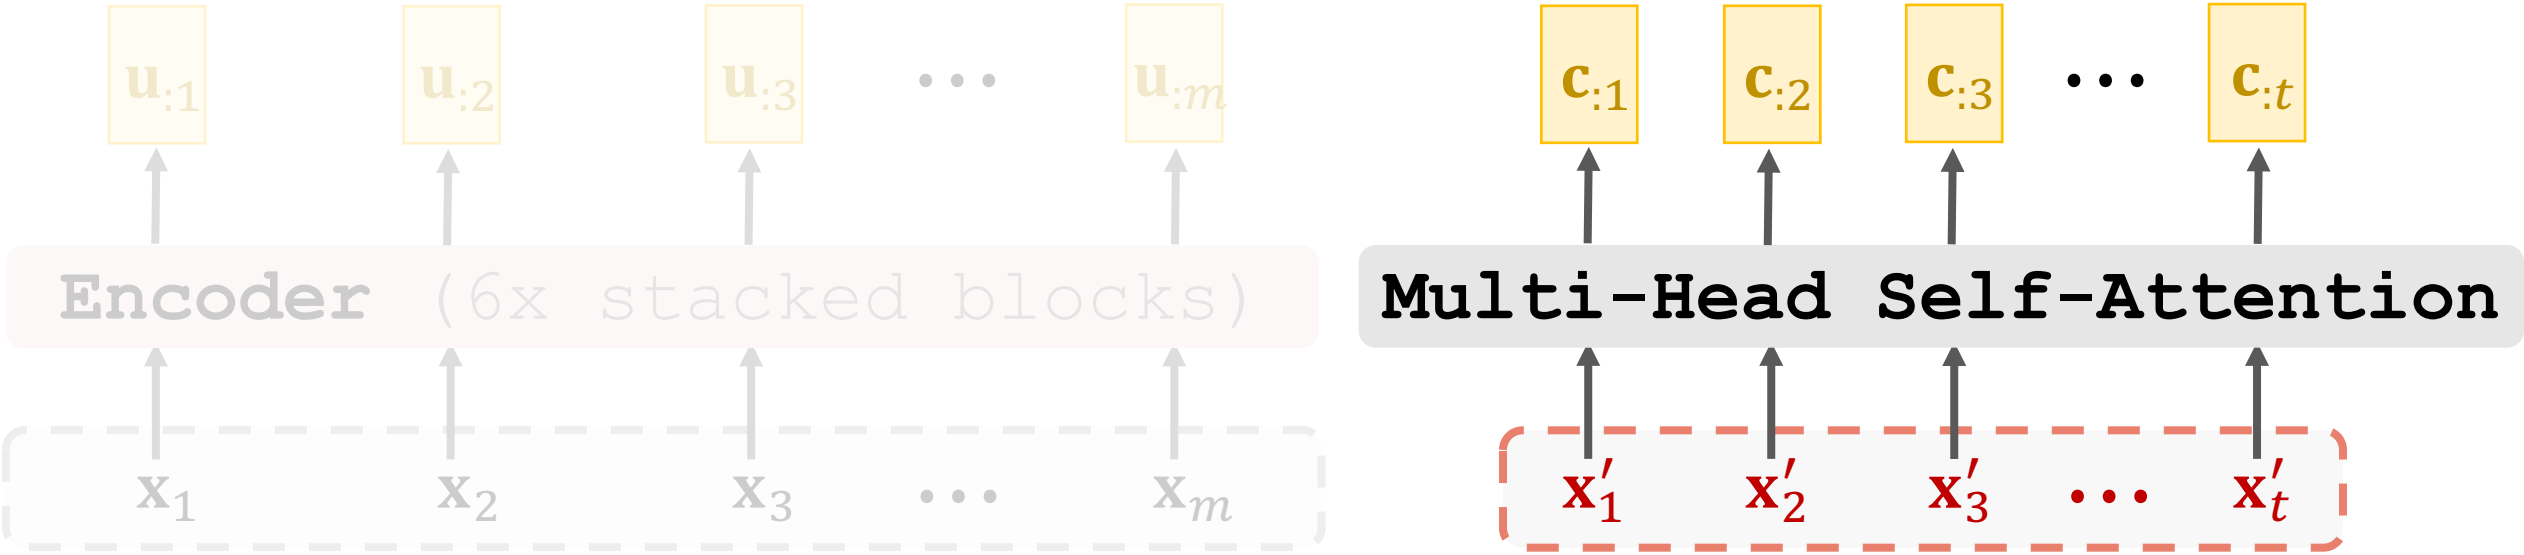

Stacked Attentions

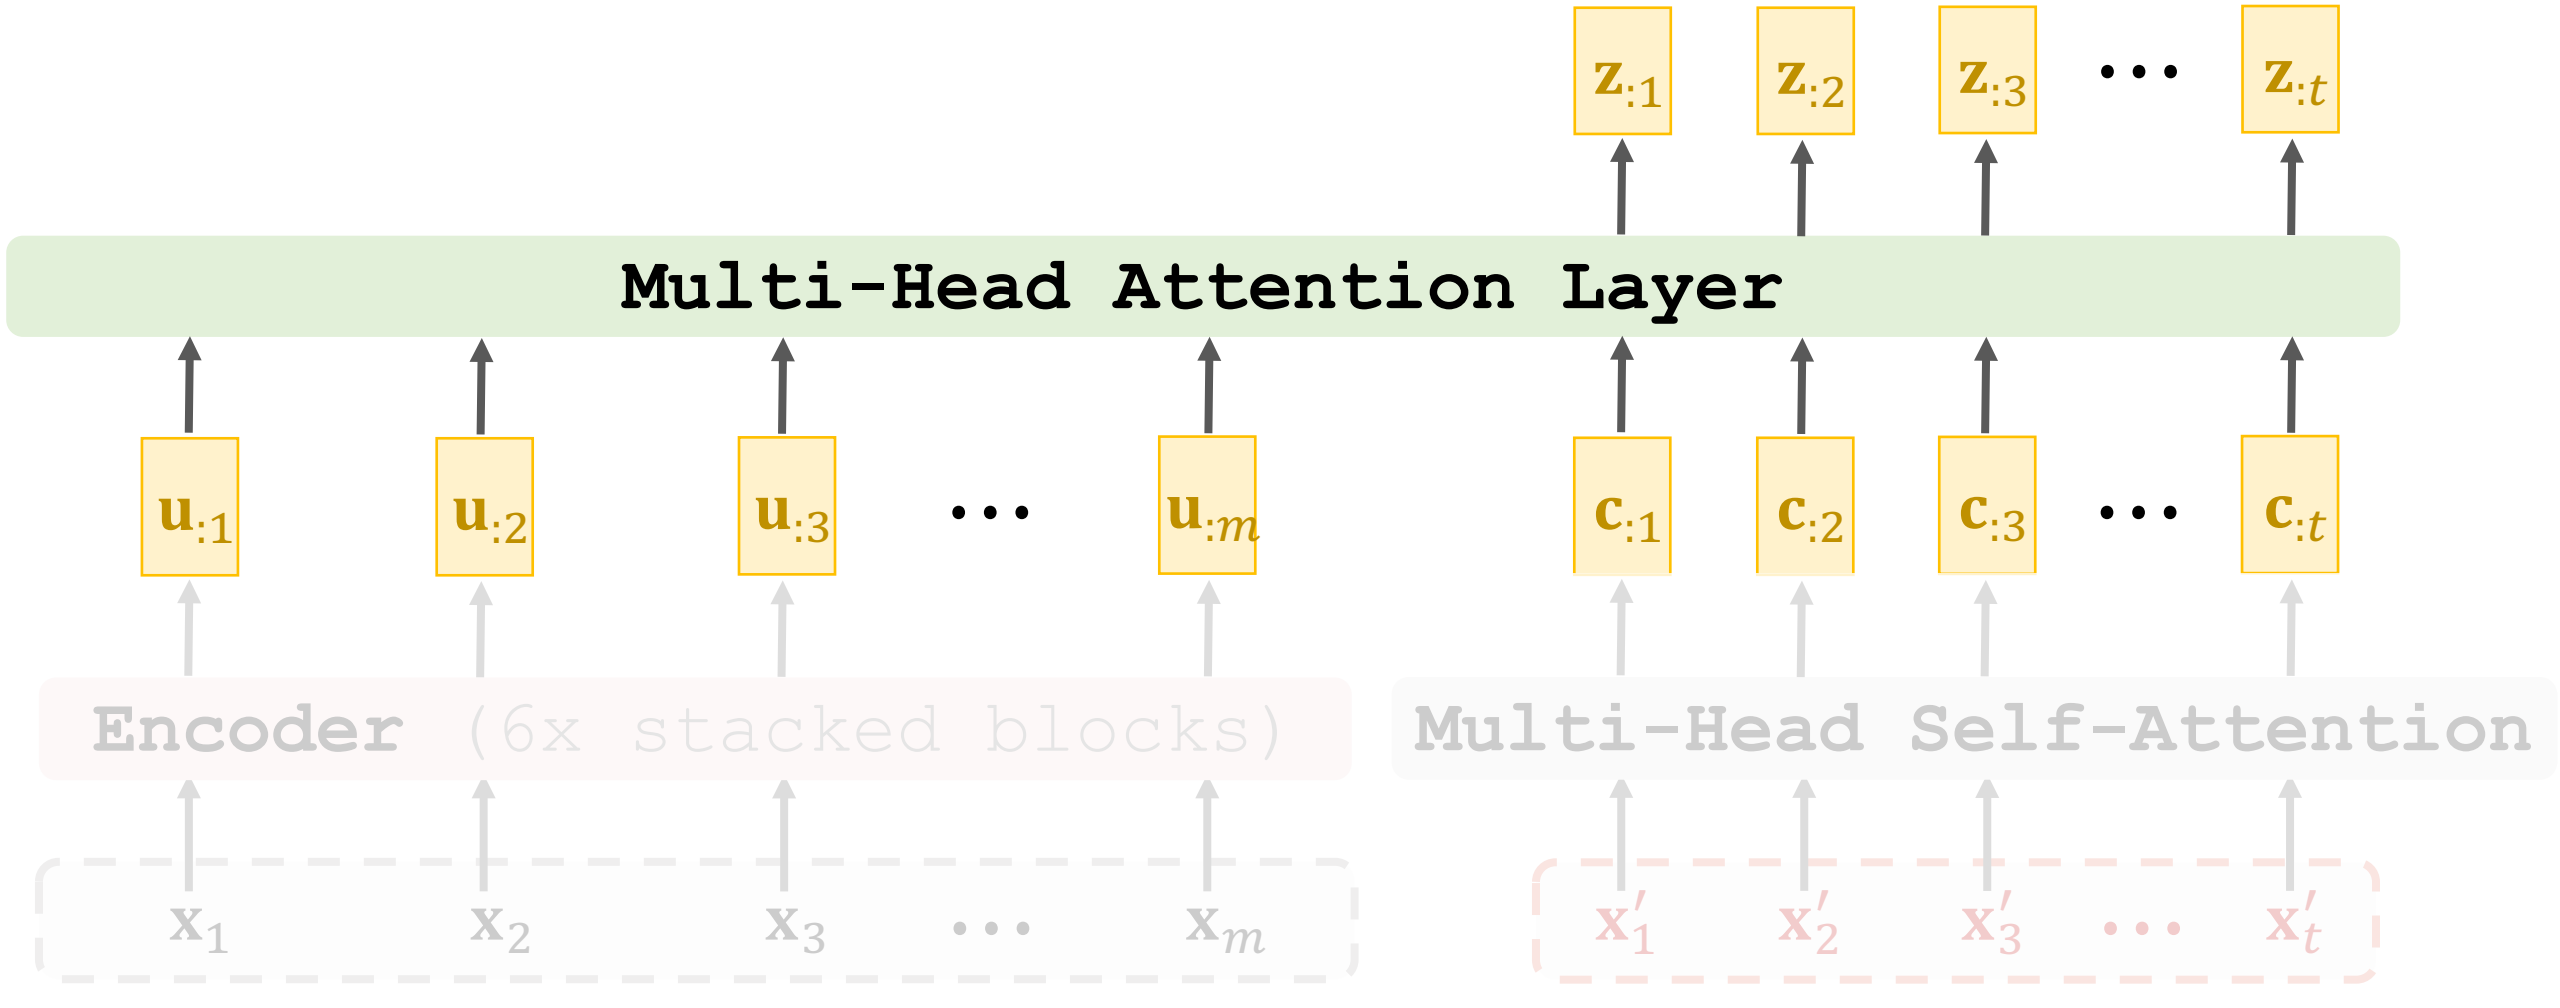

Stacked Attentions

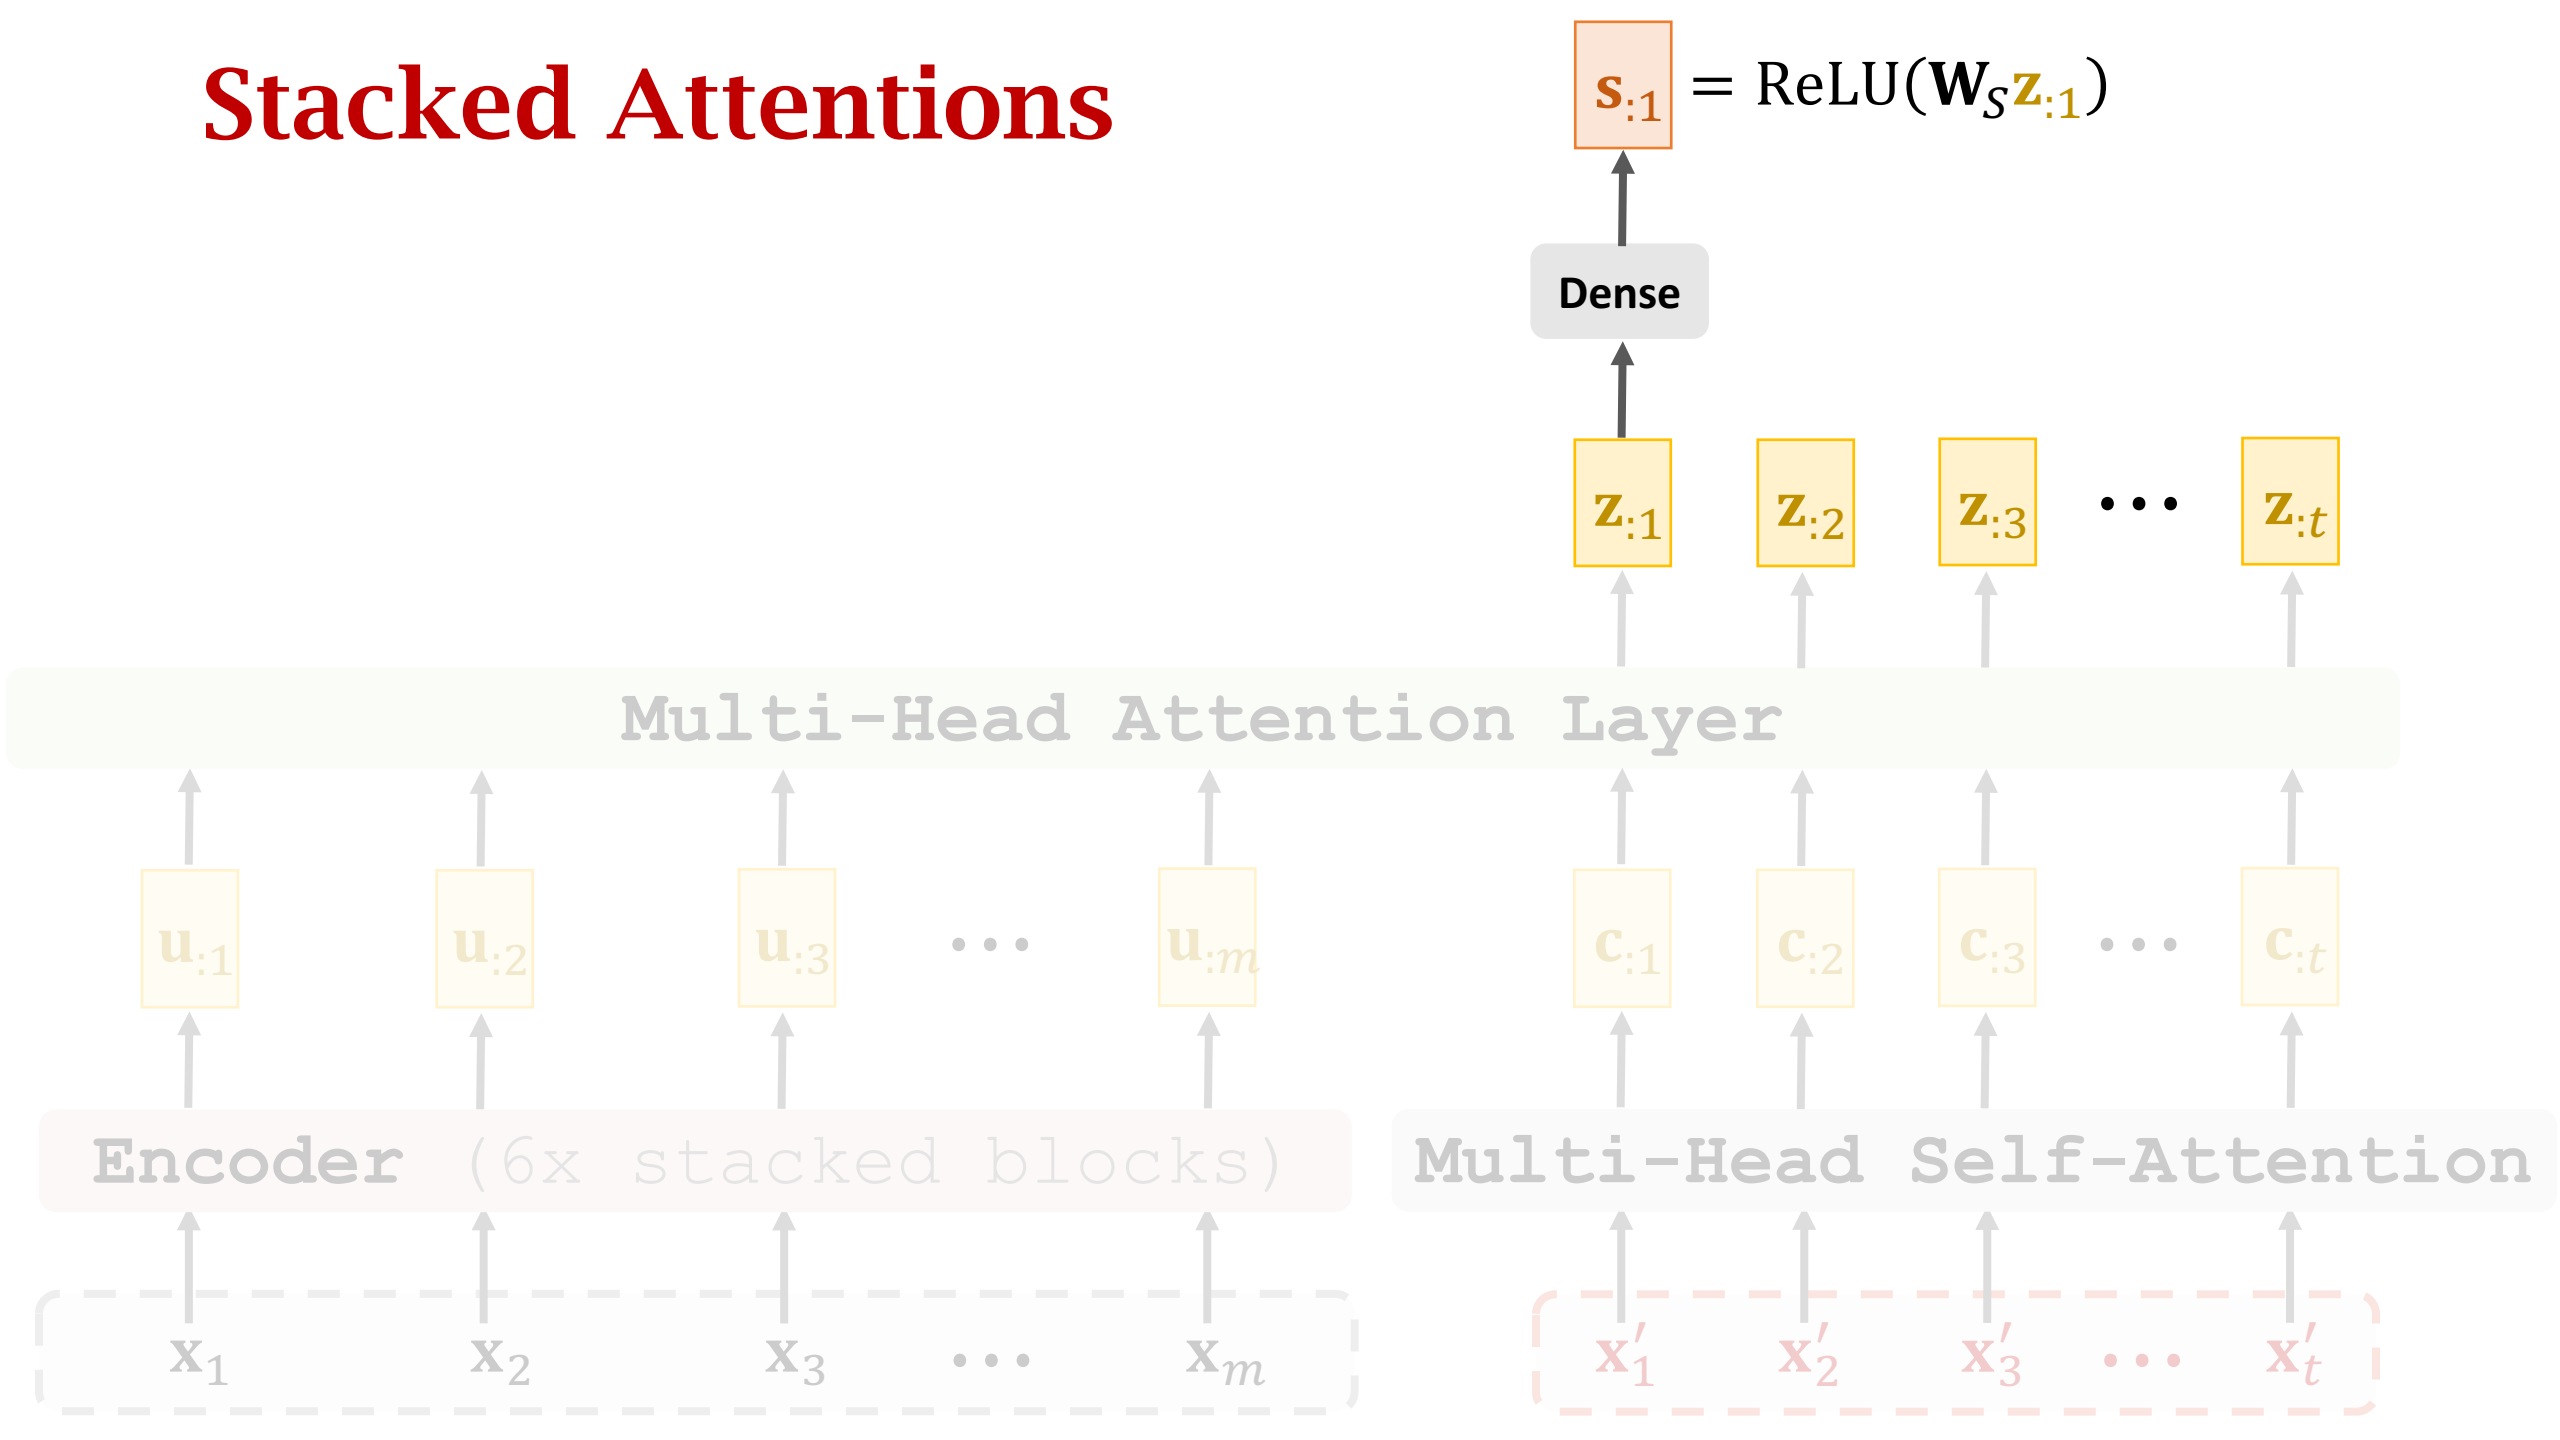

Stacked Attentions

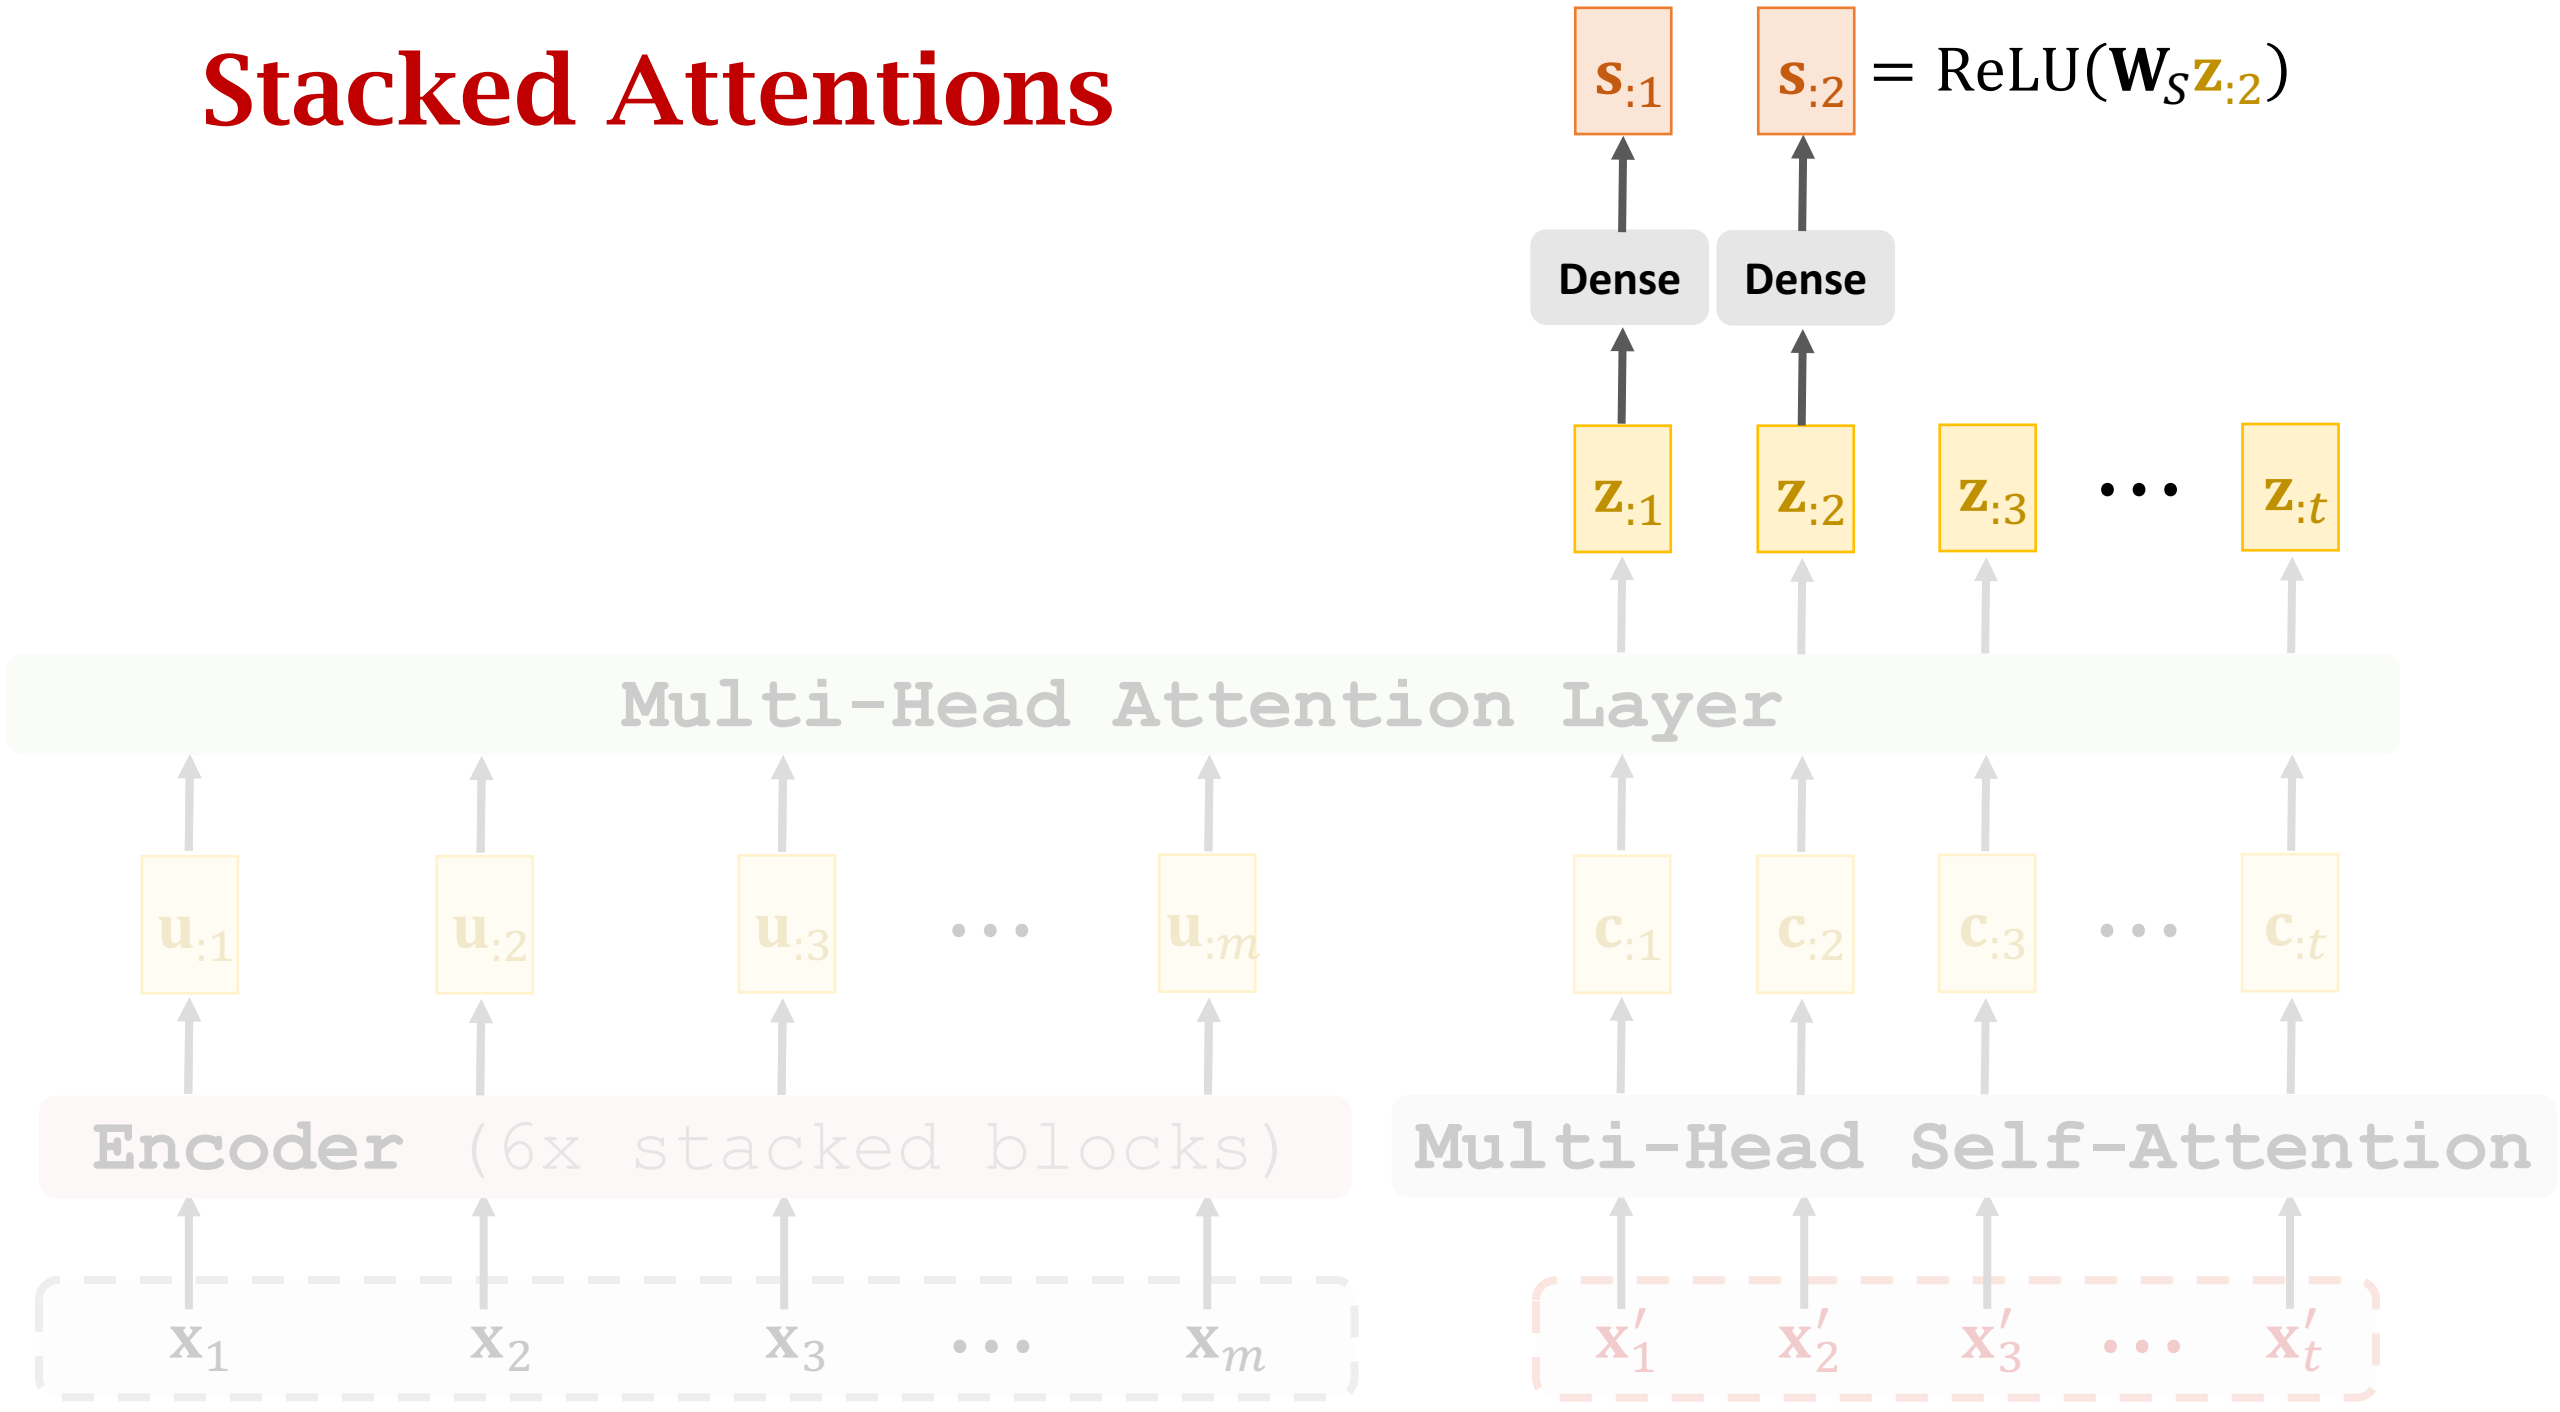

Stacked Attentions

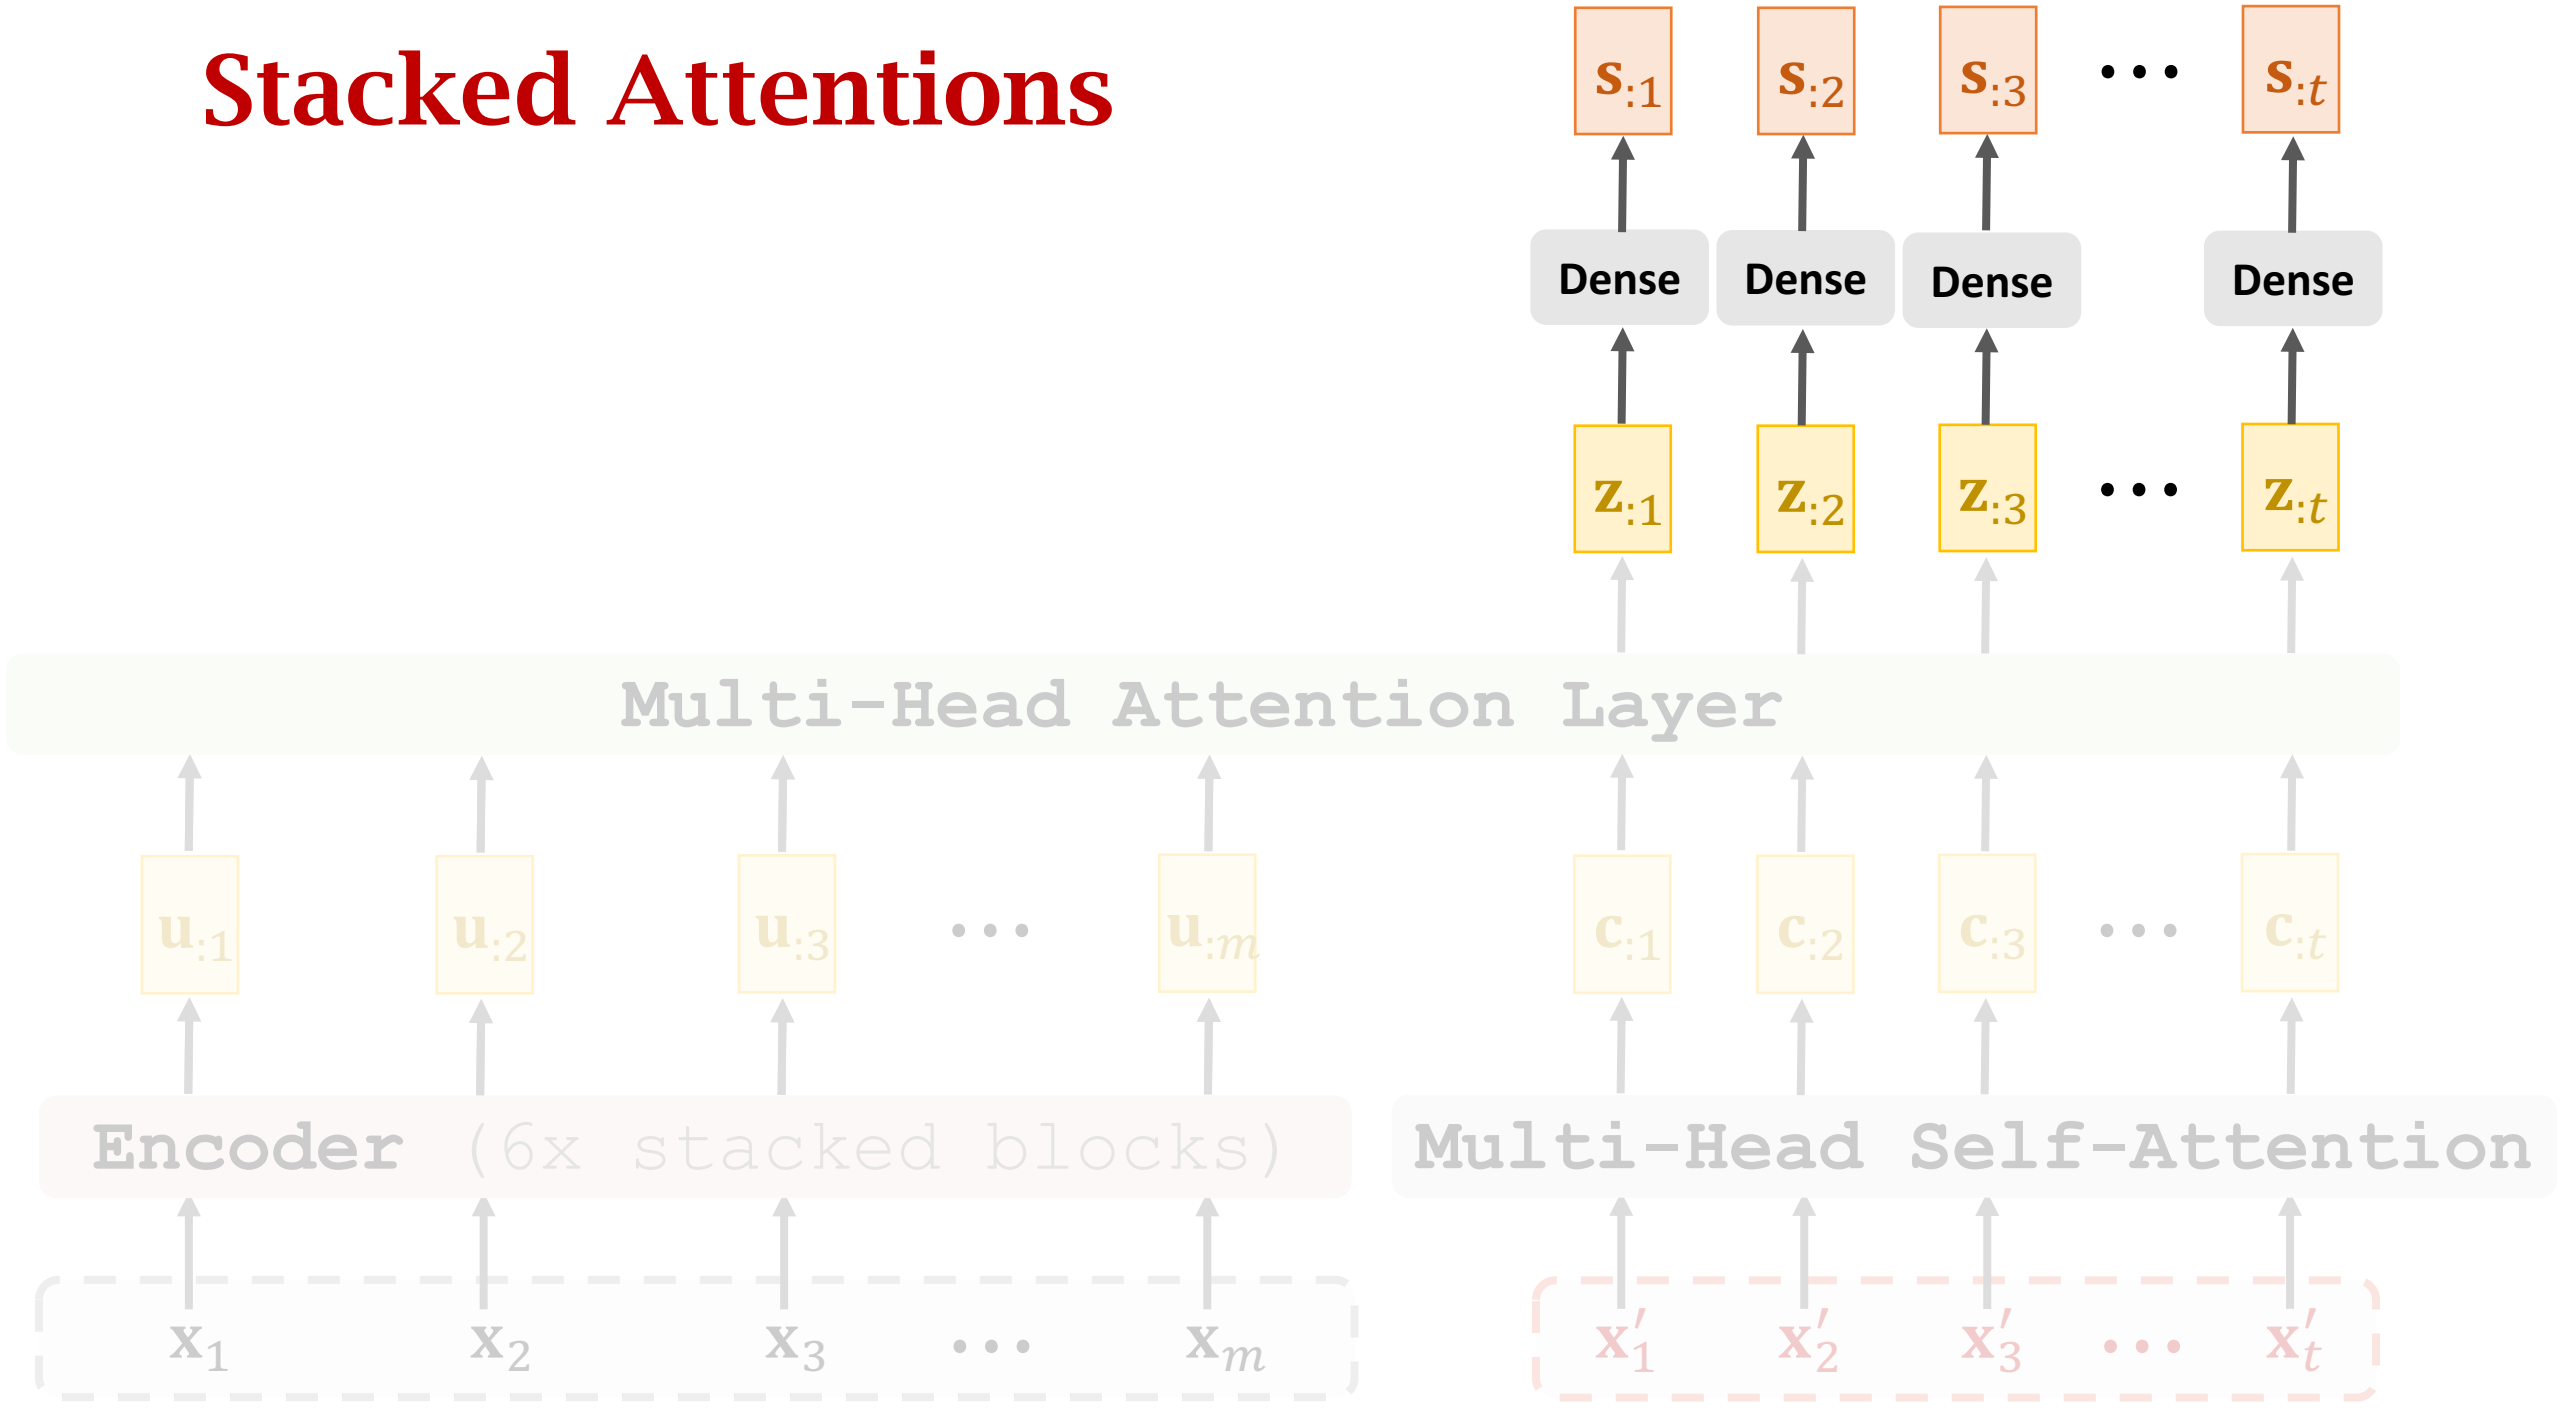

Transformer's Decoder : One Block

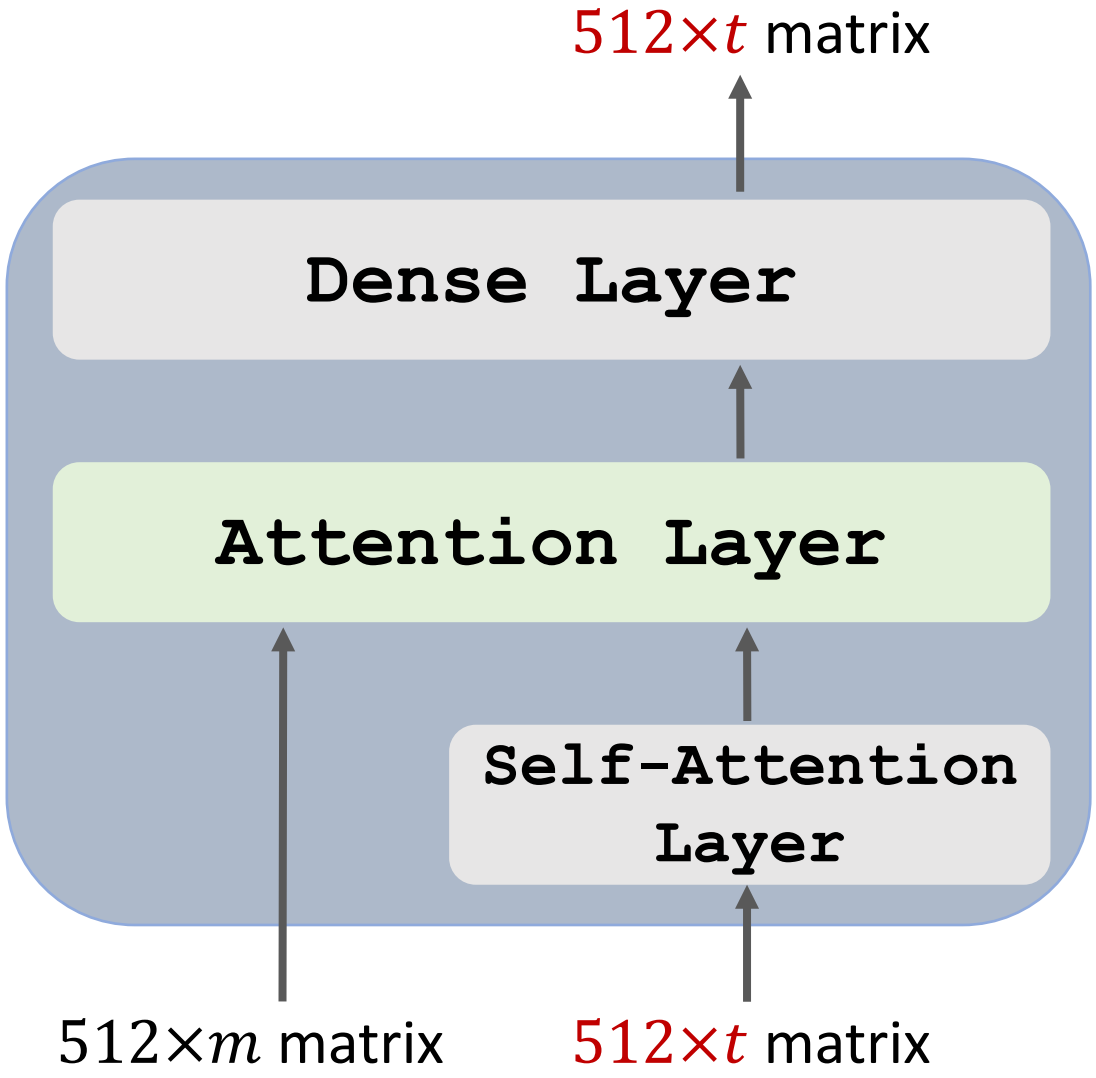

Transformer's Decoder : One Block

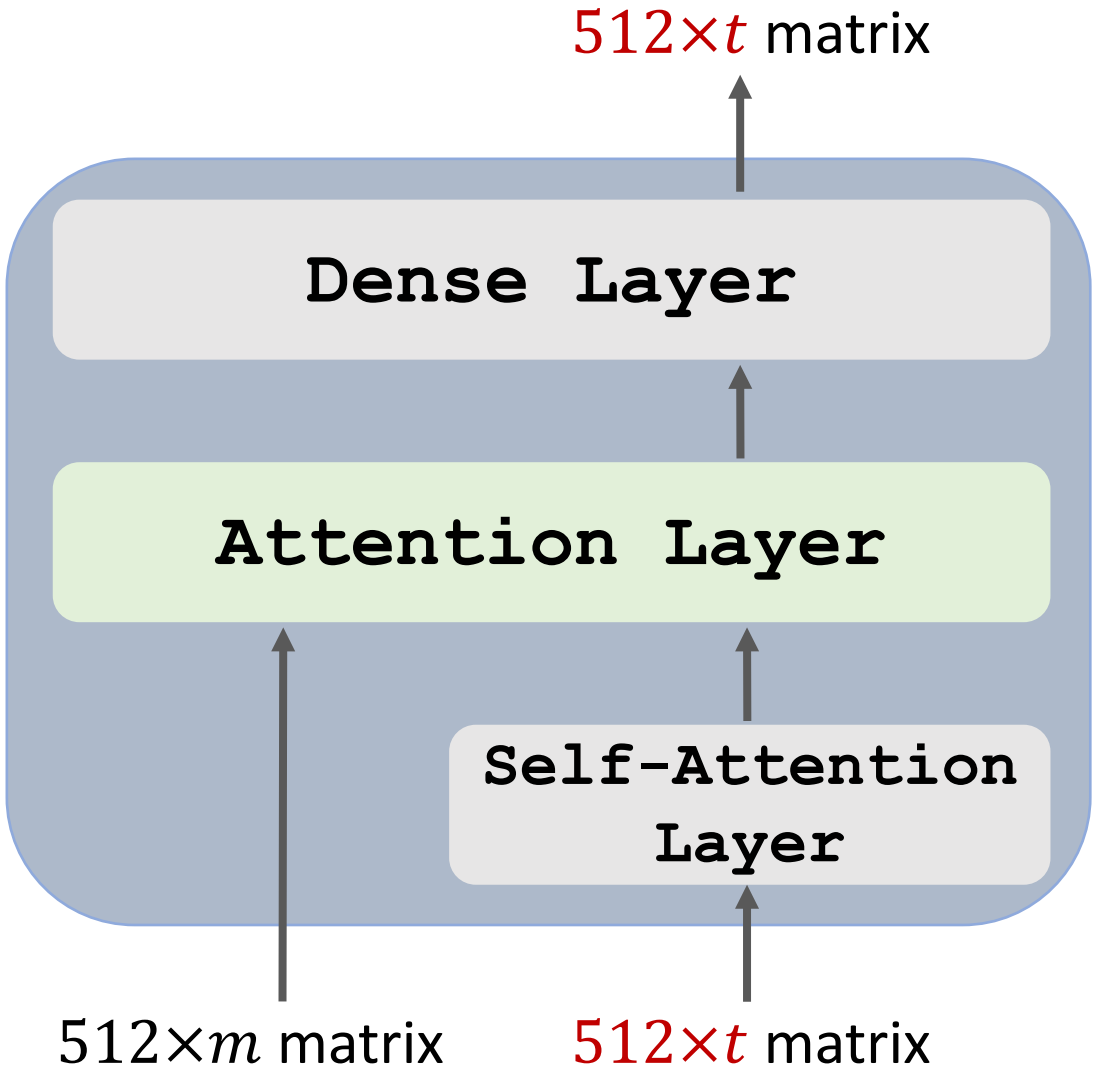

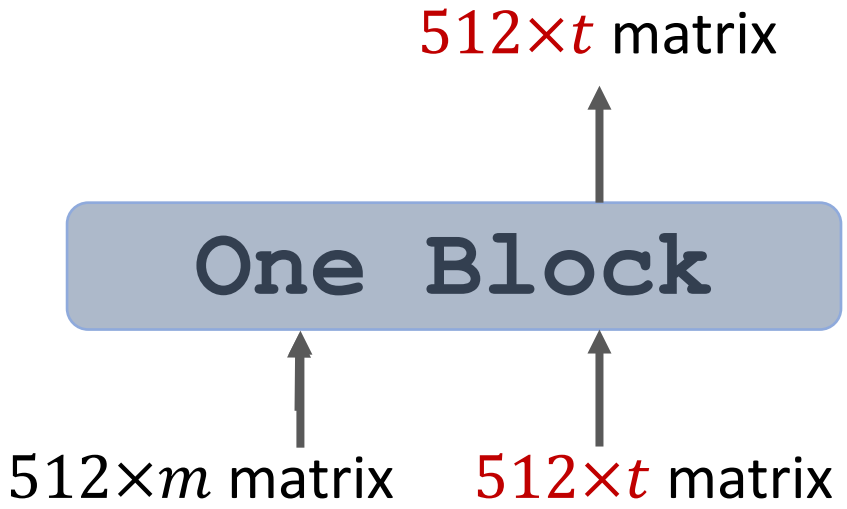

Put Everything Together

Transformer

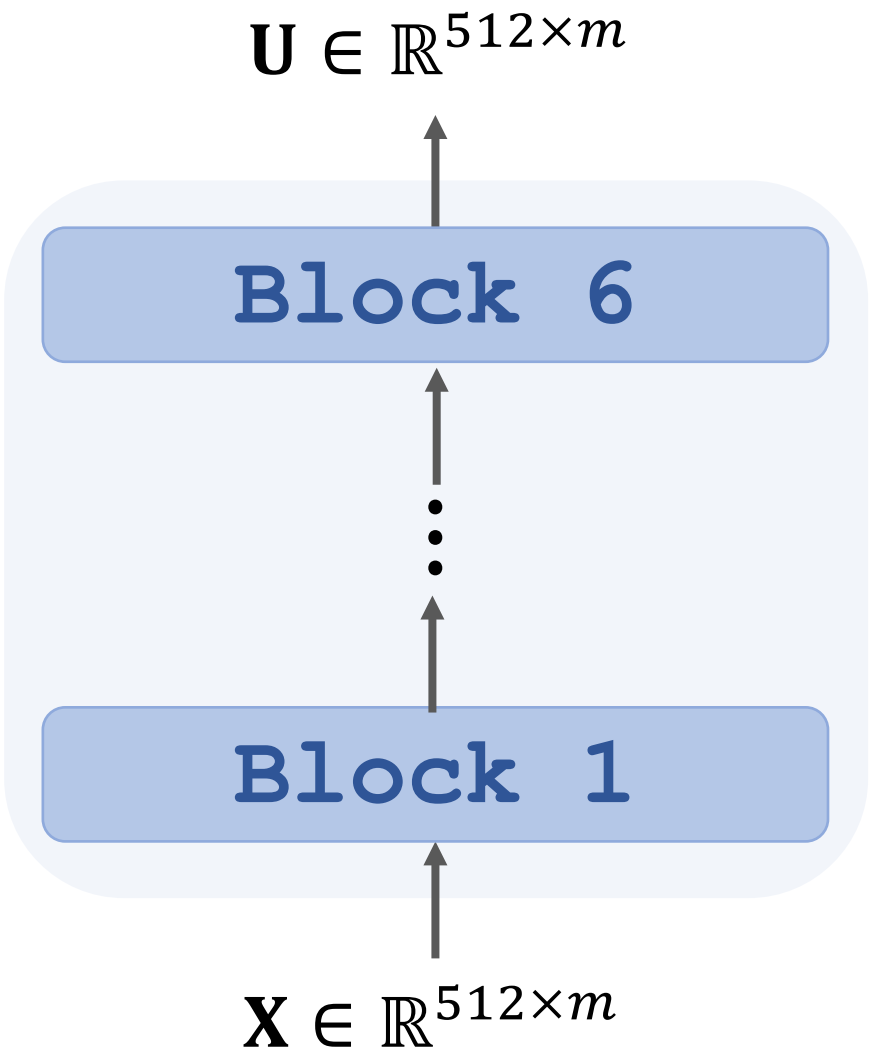

Encoder Network

Transformer

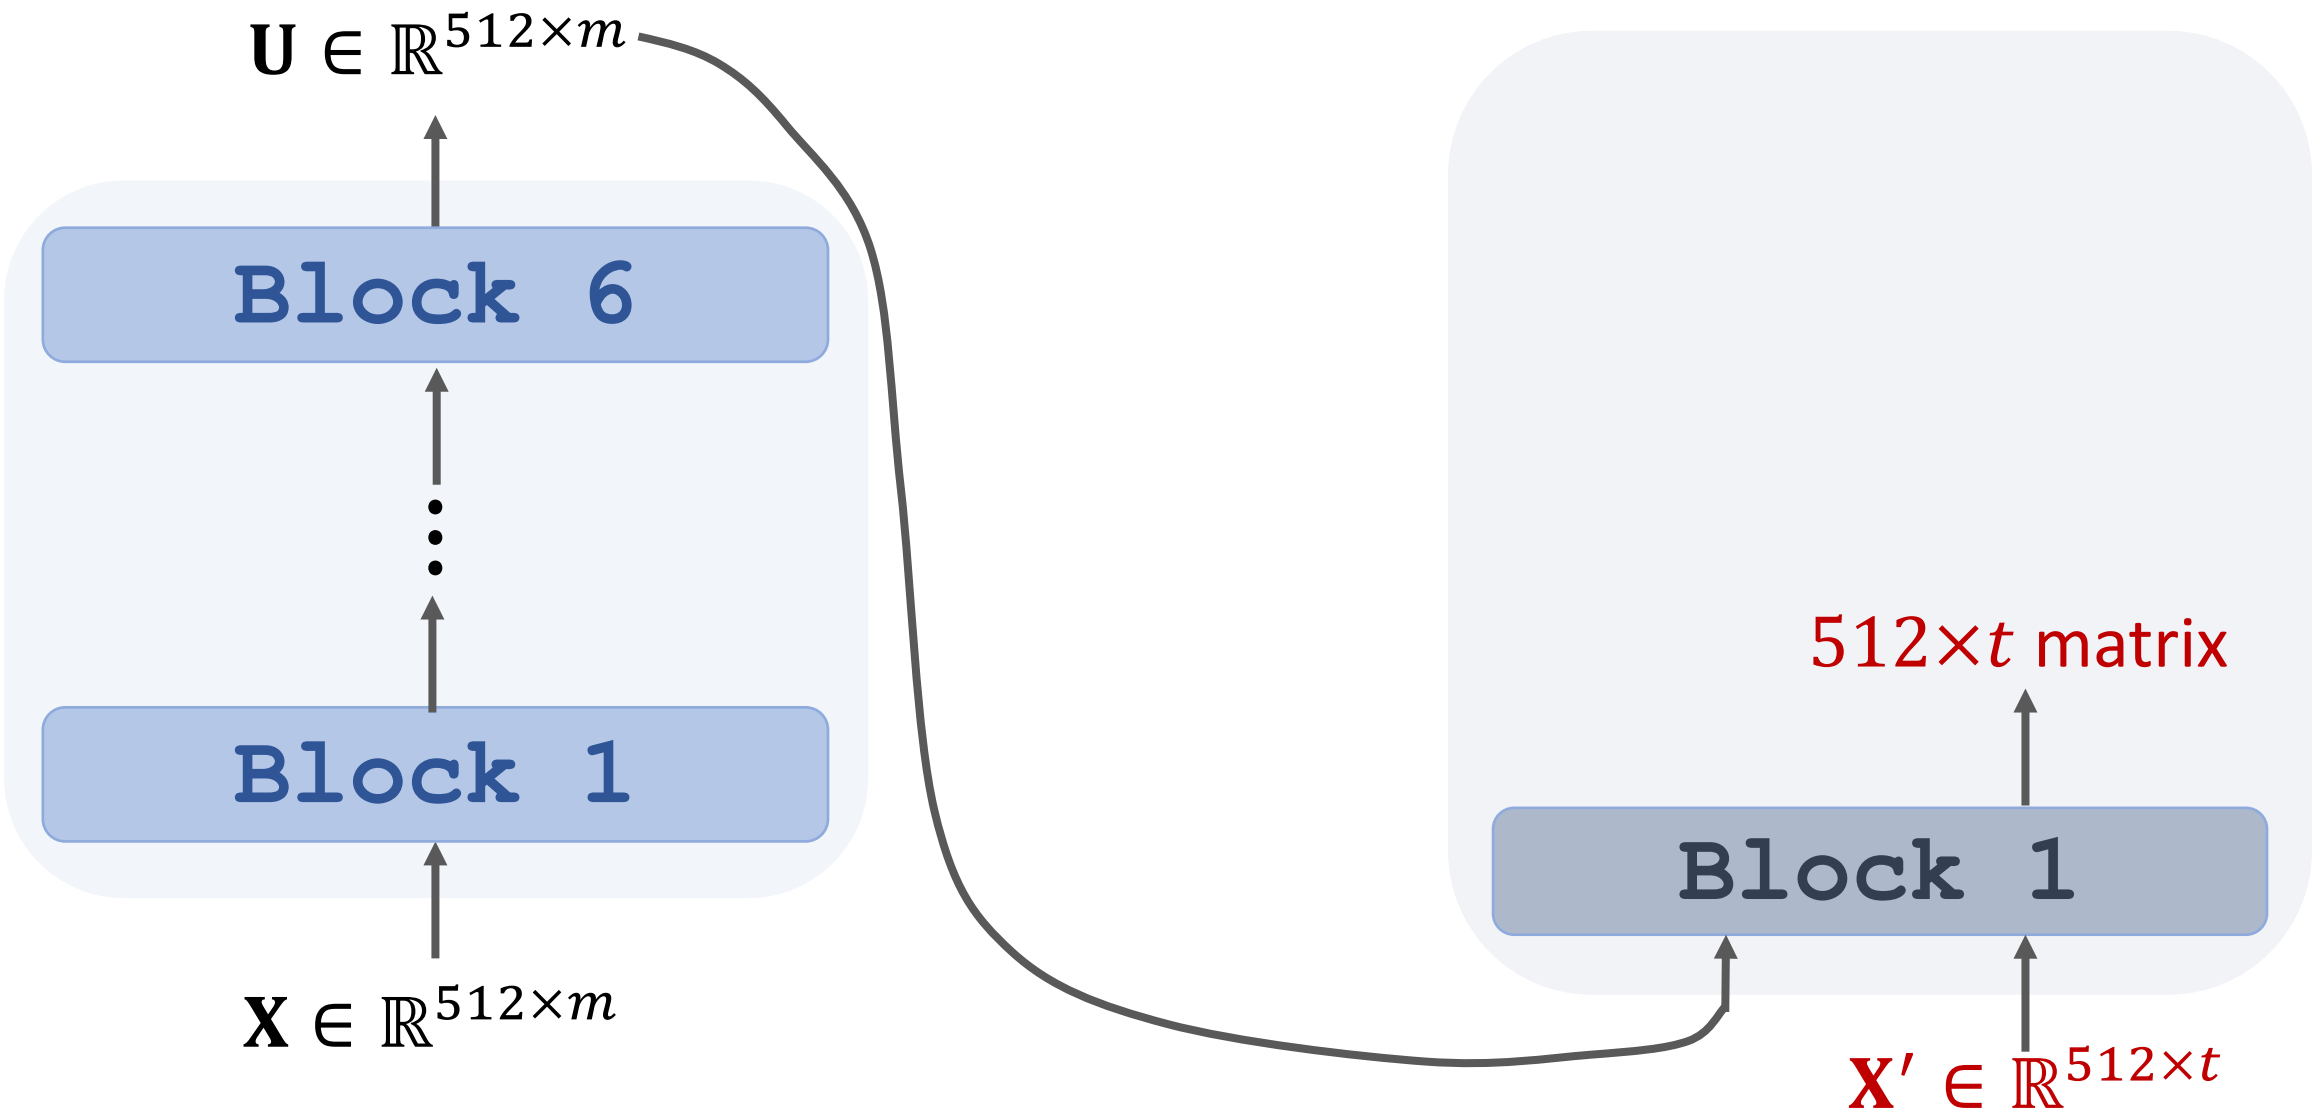

Transformer

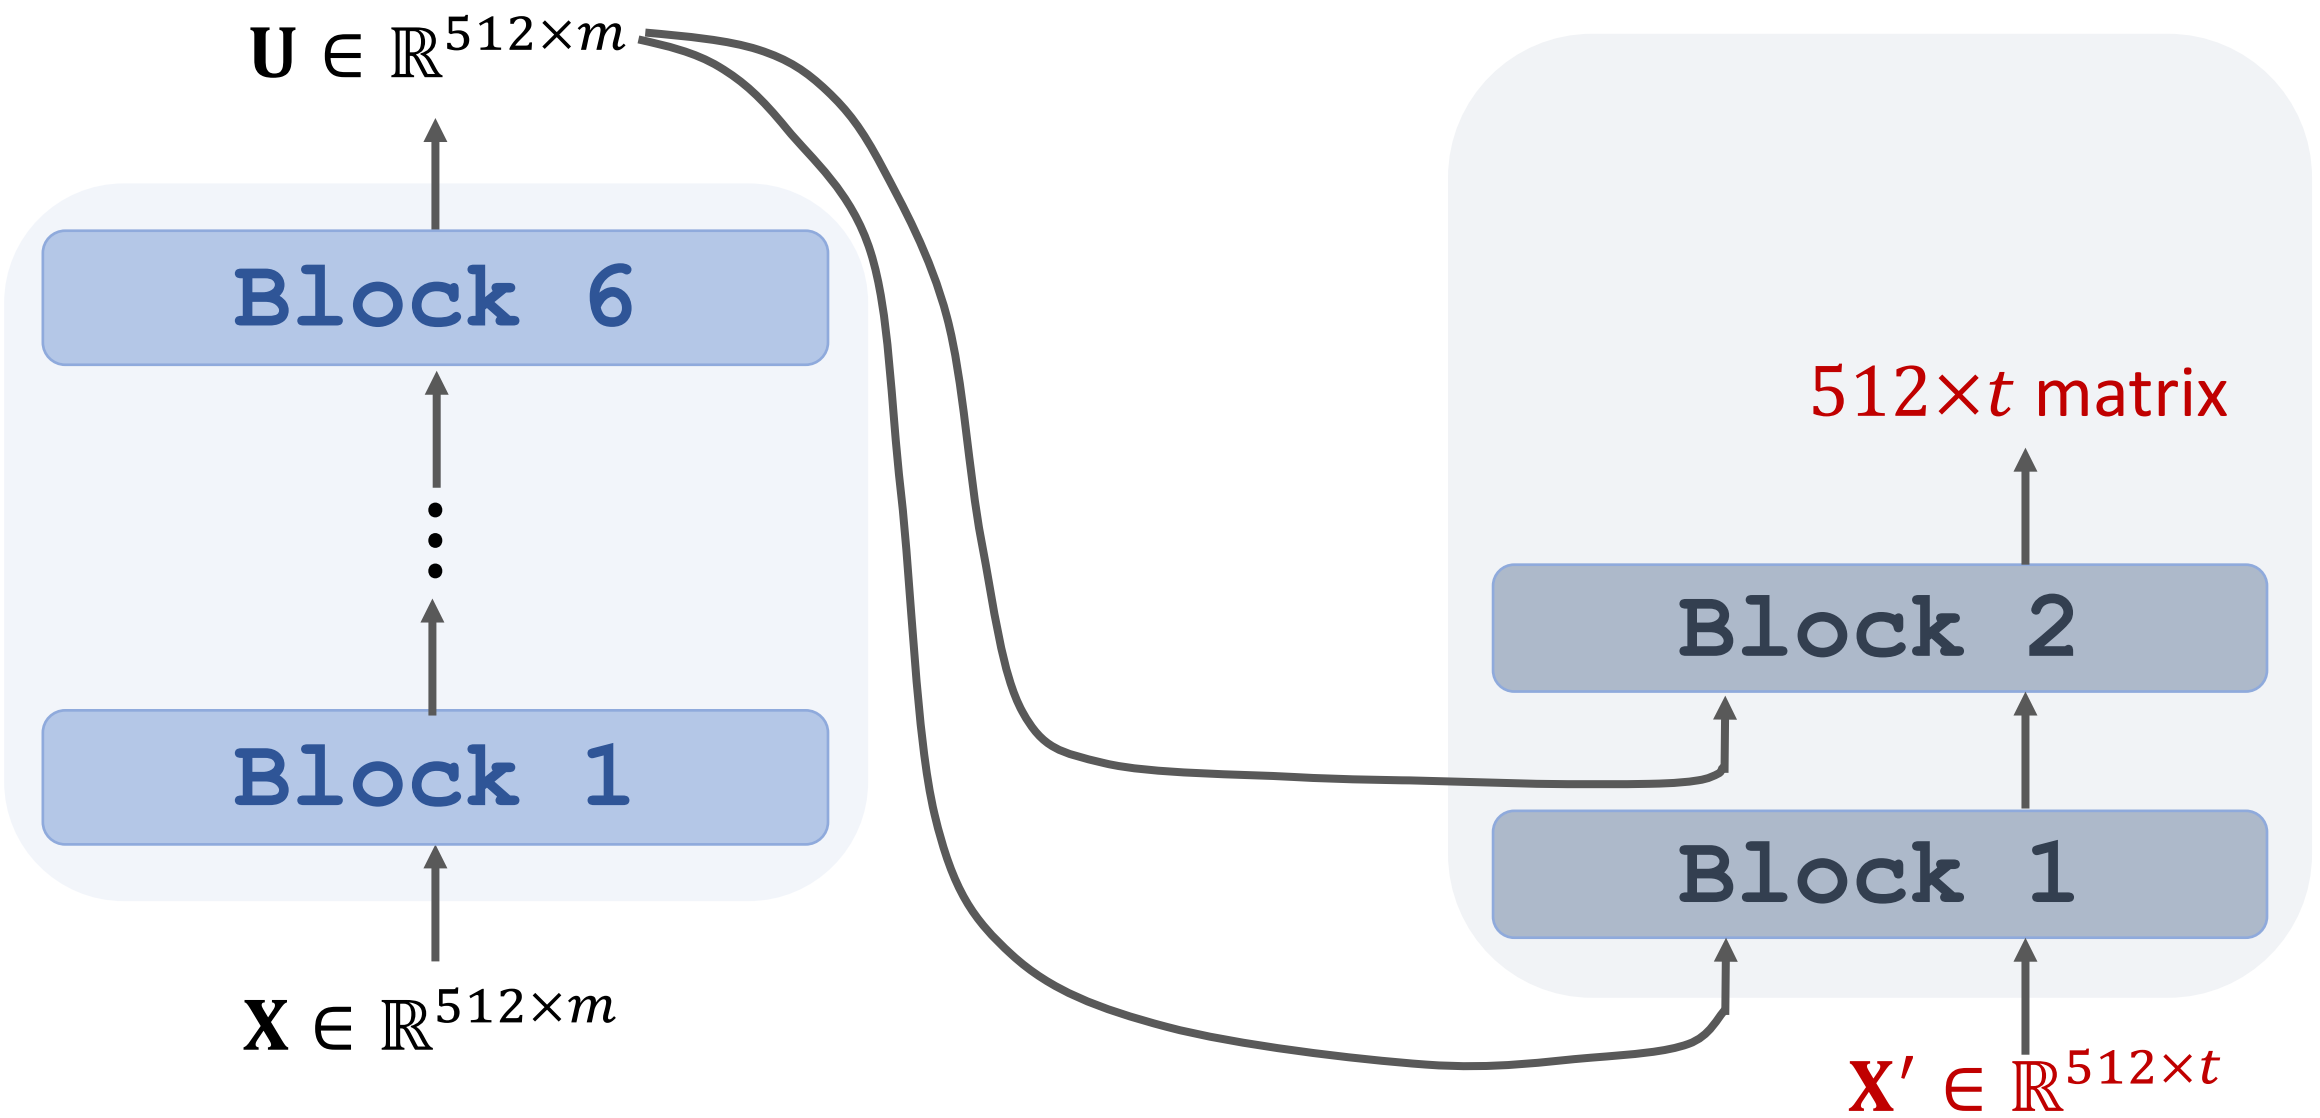

Transformer

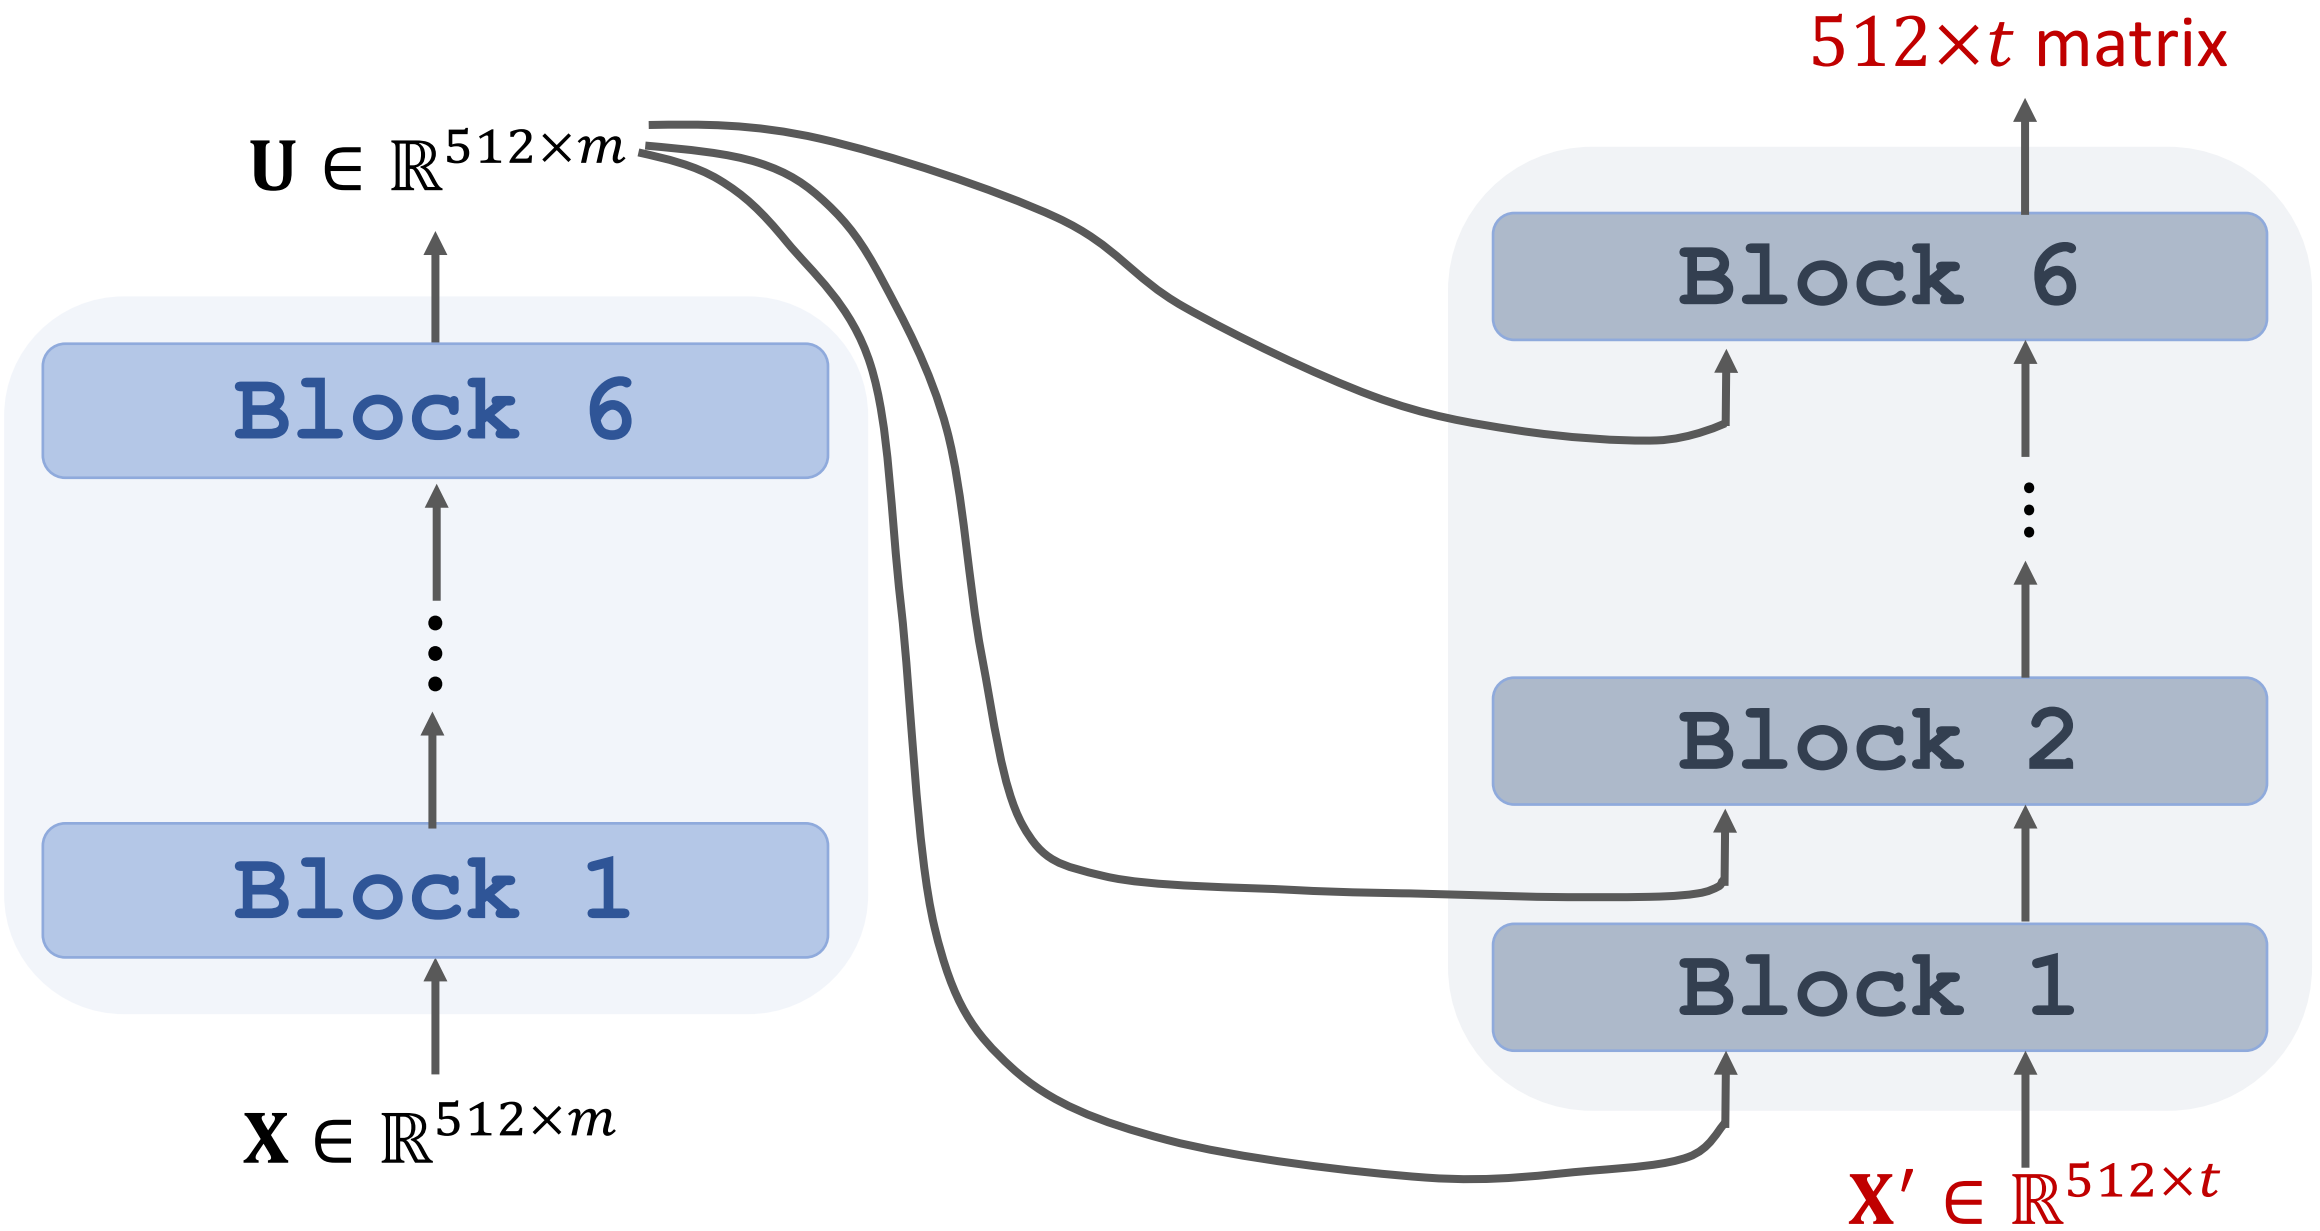

Comparison with RNN Seq2Seq Model

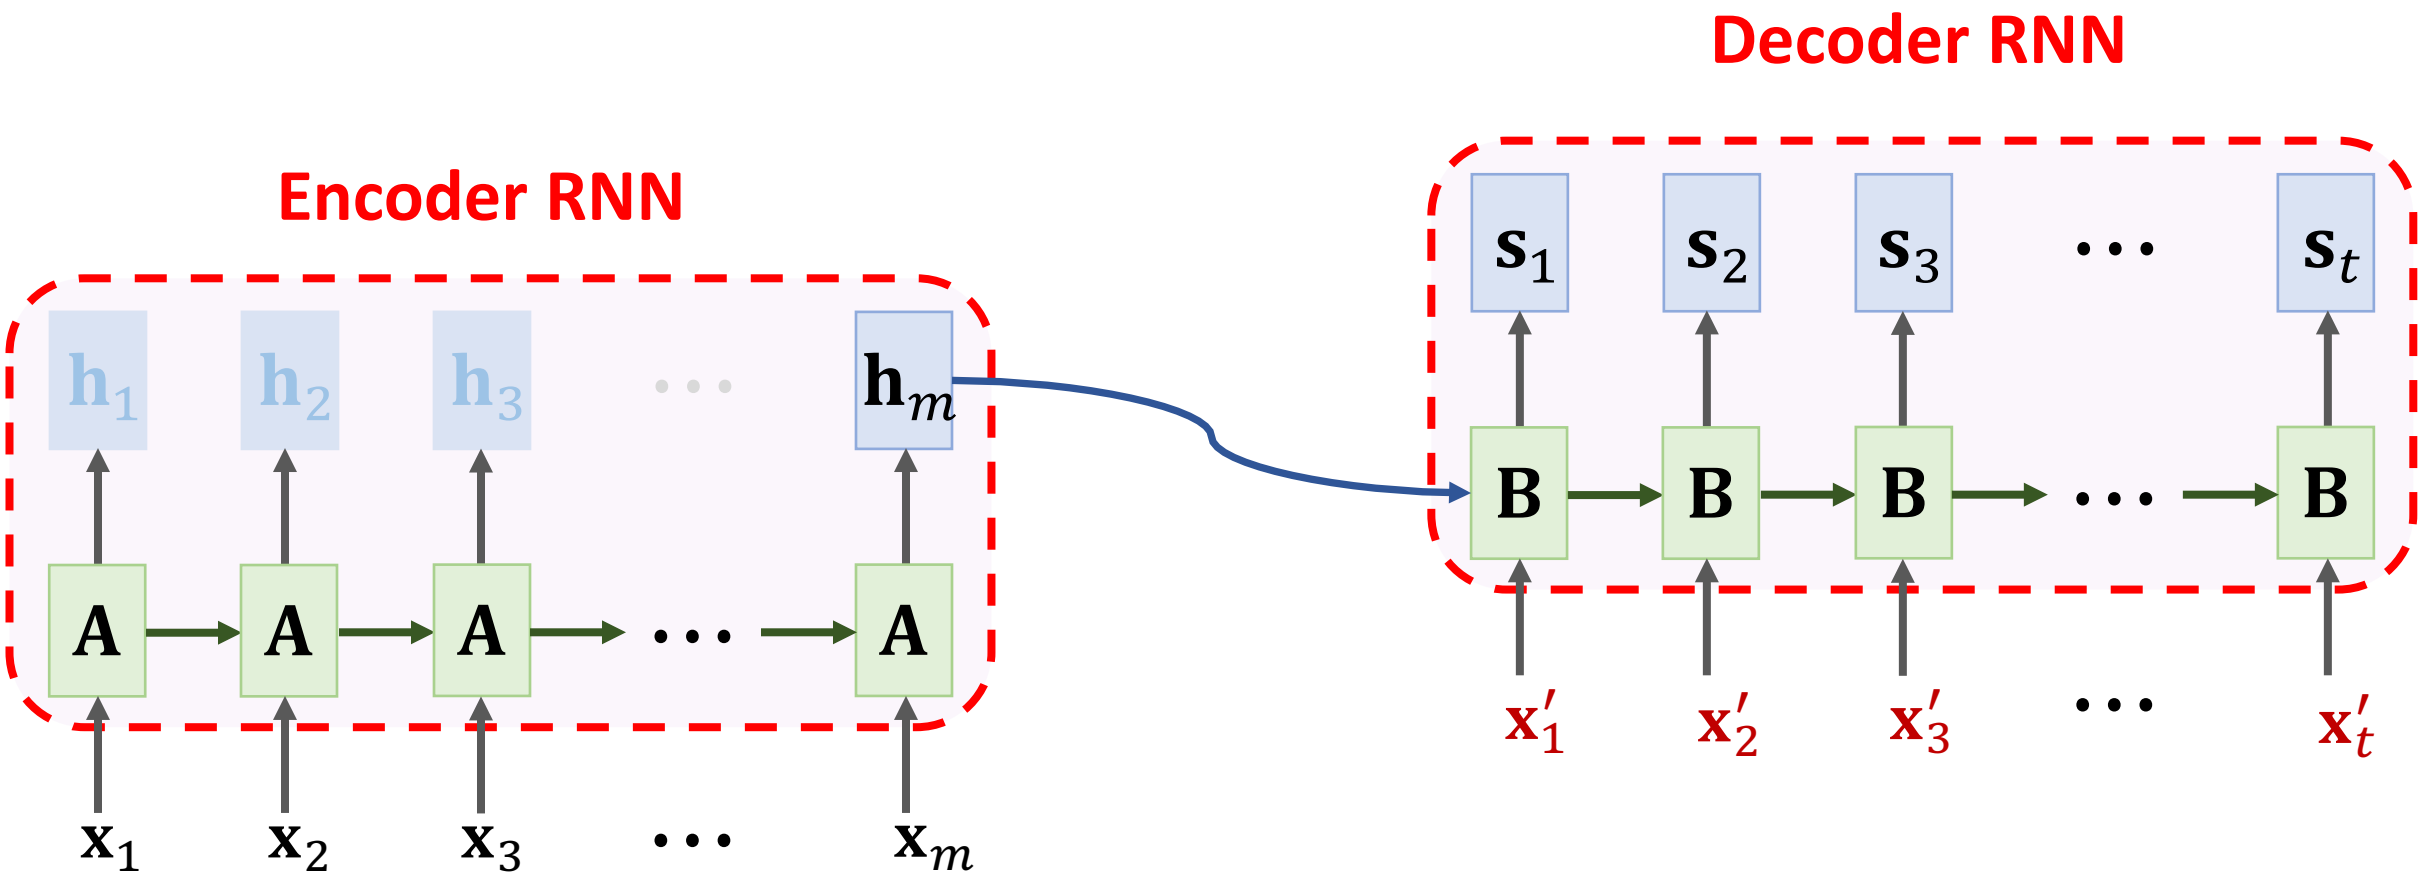

Summary

From Single-Head to Multi-Head

- Single-head self-attention can be combined to form a multi-head self-attention.

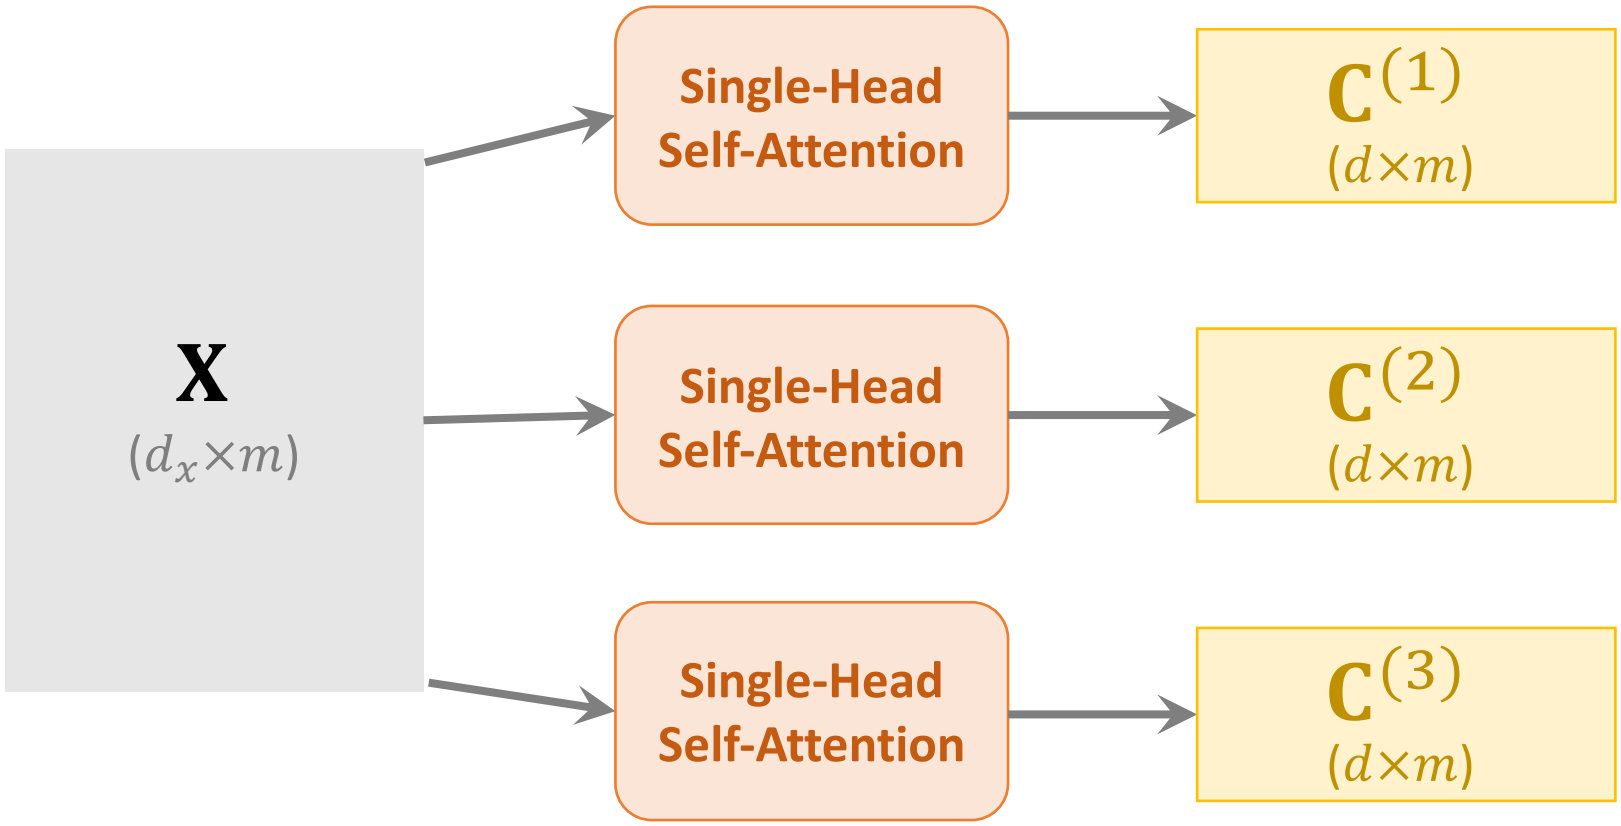

From Single-Head to Multi-Head

- Single-head self-attention can be combined to form a multi-head self-attention.

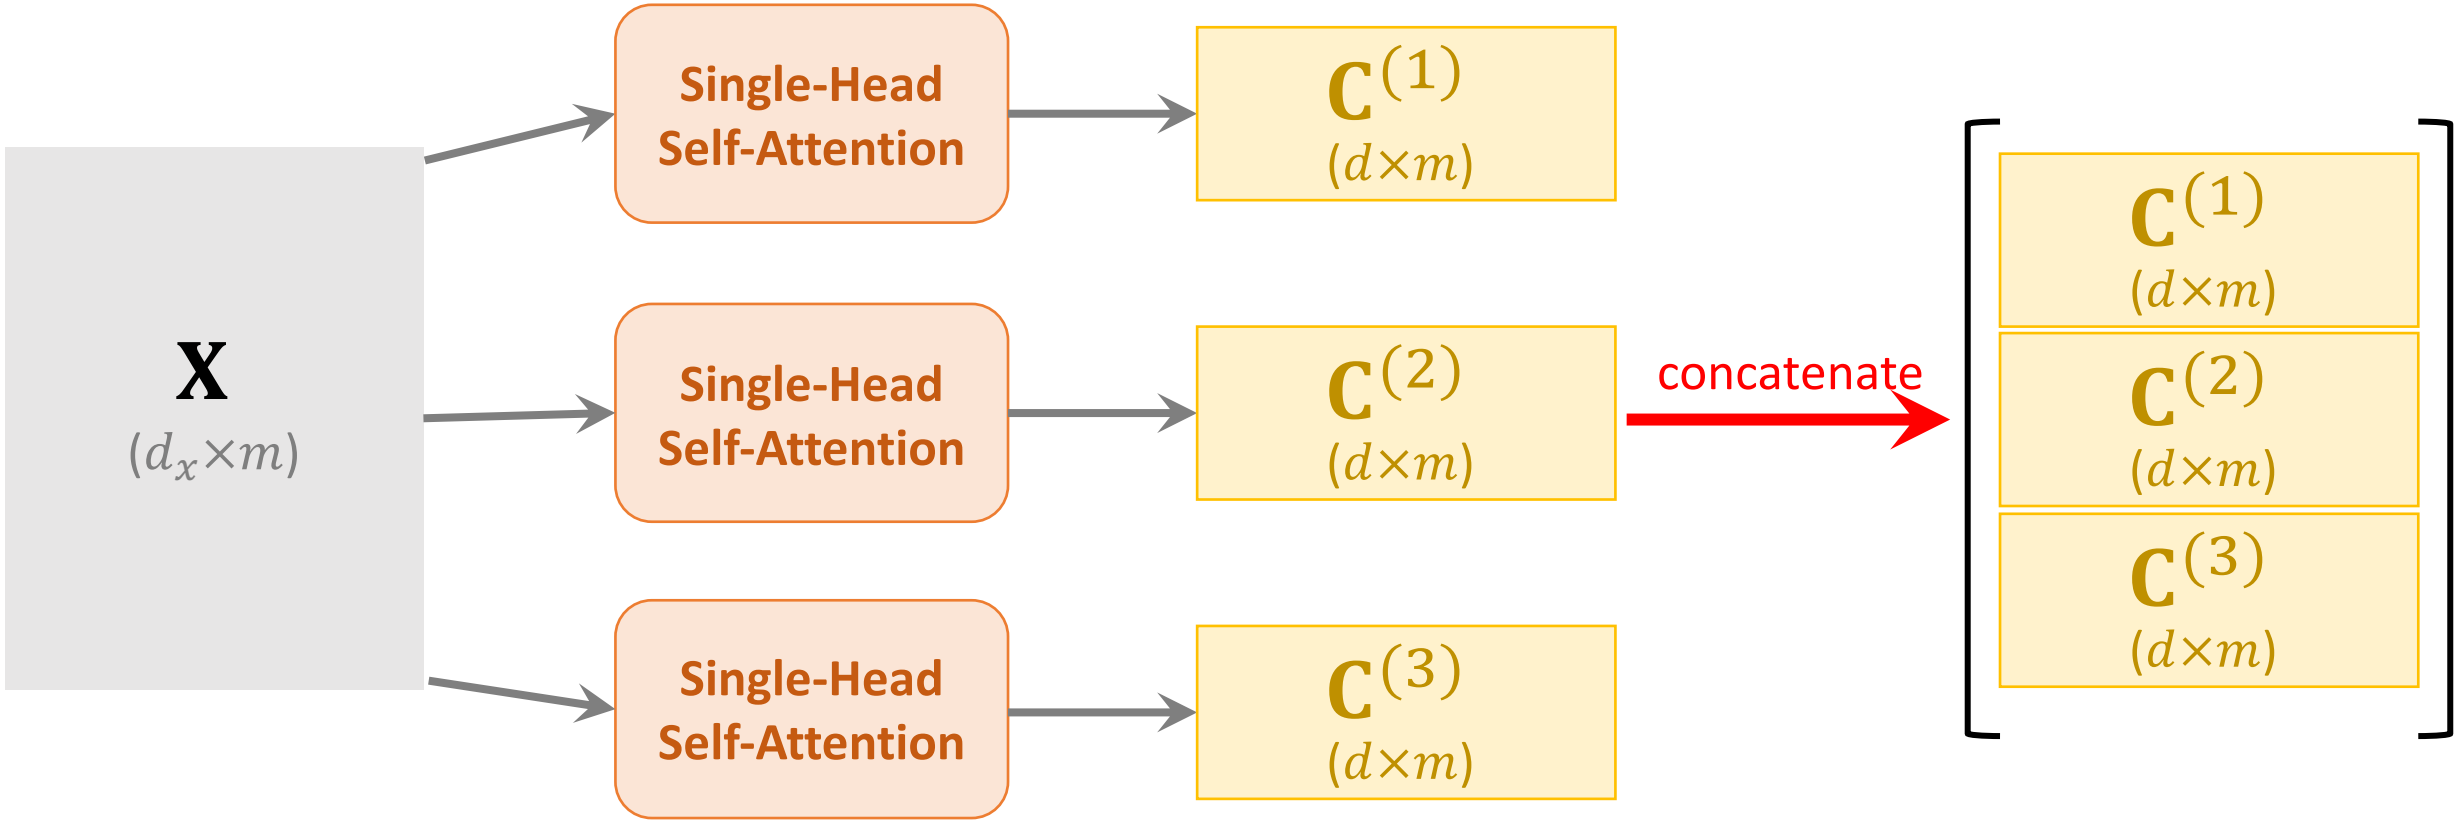

From Single-Head to Multi-Head

- Single-head self-attention can be combined to form a multi-head self-attention.

- Single-head attention can be combined to form a multi-head attention.

Encoder Network of Transformer

- 1 encoder block $\approx$ multi-head self-attention + dense.
- Input shape: $512 \times m$.
- Output shape: $512 \times m$.

- Encoder network is a stack of 6 such blocks.

Decoder Network of Transformer

- 1 decoder block $\approx$ multi-head self-attention + multi-head attention + dense.
- Input shape: $(512\times m\ ,\ 512\times t)$.
- Output shape: $512\times t$.

- Decoder network is a stack of 6 such blocks.

Transformer Model

- Transformer is Seq2Seq model; it has an encoder and a decoder.
- Transformer model is not RNN.
- Transformer is based on attention and self-attention.
- Transformer outperforms all the state-of-the-art RNN models.

Thank you!